

A Swiss Army Infinitesimal Jackknife

Ryan Giordano Will Stephenson
rgiordano@berkeley.edu wtstephe@mit.edu

Runjing Liu Michael I. Jordan
runjing.liu@berkeley.edu jordan@cs.berkeley.edu

Tamara Broderick
tbroderick@csail.mit.edu

December 16, 2019

Abstract

The error or variability of machine learning algorithms is often assessed by repeatedly re-fitting a model with different weighted versions of the observed data. The ubiquitous tools of cross-validation (CV) and the bootstrap are examples of this technique. These methods are powerful in large part due to their model agnosticism but can be slow to run on modern, large data sets due to the need to repeatedly re-fit the model. In this work, we use a linear approximation to the dependence of the fitting procedure on the weights, producing results that can be faster than repeated re-fitting by an order of magnitude. This linear approximation is sometimes known as the “infinitesimal jackknife” in the statistics literature, where it is mostly used as a theoretical tool to prove asymptotic results. We provide explicit finite-sample error bounds for the infinitesimal jackknife in terms of a small number of simple, verifiable assumptions. Our results apply whether the weights and data are stochastic or deterministic, and so can be used as a tool for proving the accuracy of the infinitesimal jackknife on a wide variety of problems. As a corollary, we state mild regularity conditions under which our approximation consistently estimates true leave- k -out cross-validation for any fixed k . These theoretical results, together with modern automatic differentiation software, support the application of the infinitesimal jackknife to a wide variety of practical problems in machine learning, providing a “Swiss Army infinitesimal jackknife.” We demonstrate the accuracy of our methods on a range of simulated and real datasets.

1 Introduction

Statistical machine learning methods are increasingly deployed in real-world problem domains where they are the basis of decisions affecting individuals’

employment, savings, health, and safety. Unavoidable randomness in data collection necessitates understanding how our estimates, and resulting decisions, might have differed had we observed different data. Both cross validation (CV) and the bootstrap attempt to diagnose this variation and are widely used in classical data analysis. But these methods are often prohibitively slow for modern, massive datasets, as they require running a learning algorithm on many slightly different datasets. In this work, we propose to replace these many runs with a single perturbative approximation. We show that the computation of this approximation is far cheaper than the classical methods, and we provide theoretical conditions that establish its accuracy.

Many data analyses proceed by minimizing a loss function of exchangeable data. Examples include empirical loss minimization and M-estimation based on product likelihoods. Since we typically do not know the true distribution generating the data, it is common to approximate the dependence of our estimator on the data via the dependence of the estimator on the empirical distribution. In particular, we often form a new, proxy dataset using random or deterministic modifications of the empirical distribution, such as randomly removing k datapoints for leave- k -out CV. A proxy dataset obtained in this way can be represented as a weighting of the original data. From a set of such proxy datasets we can obtain estimates of uncertainty, including estimates of bias, variance, and prediction accuracy.

As data and models grow, the cost of repeatedly solving a large optimization problem for a number of different values of weights can become impractically large. Conversely, though, larger datasets often exhibit greater regularity; in particular, under fairly general conditions, limit laws based on independence imply that an optimum exhibits diminishing dependence on any fixed set of data points. We use this observation to derive a linear approximation to resampling that needs to be calculated only once, but which nonetheless captures the variability inherent in the repeated computations of classical CV. Our method is an instance of the *infinitesimal jackknife* (IJ), a general methodology that was historically a precursor to cross-validation and the bootstrap [Jaeckel, 1972, Efron, 1982]. Part of our argument is that variants of the IJ should be reconsidered for modern large-scale applications because, for smooth optimization problems, the IJ can be calculated automatically with modern automatic differentiation tools [Baydin et al., 2017].

By using this linear approximation, we incur the cost of forming and inverting a matrix of second derivatives with size equal to the dimension of the parameter space, but we avoid the cost of repeatedly re-optimizing the objective. As we demonstrate empirically, this tradeoff can be extremely favorable in many problems of interest.

Our approach aims to provide a felicitous union of two schools of thought. In statistics, the IJ is typically used to prove normality or consistency of other estimators [Fernholz, 1983, Shao, 1993, Shao and Tu, 2012]. However, the conditions that are required for these asymptotic analyses to hold are prohibitively restrictive for machine learning—specifically, they require objectives with bounded gradients. A number of recent papers in machine learning have provided related

linear approximations for the special case of leave-one-out cross-validation [Koh and Liang, 2017, Rad and Maleki, 2018, Beirami et al., 2017], though their analyses lack the generality of the statistical perspective.

We combine these two approaches by modifying the proof of the Fréchet differentiability of M-estimators developed by Clarke [1983]. Specifically, we adapt the proof away from the question of Fréchet differentiability within the class of all empirical distributions to the narrower problem of approximating the exact re-weighting on a particular dataset with a potentially restricted set of weights. This limitation of what we expect from the approximation is crucial; it allows us to bound the error in terms of a complexity measure of the set of derivatives of the observed objective function, providing a basis for non-asymptotic applications in large-scale machine learning, even for objectives with unbounded derivatives. Together with modern automatic differentiation tools, these results extend the use of the IJ to a wider range of practical problems. Thus, our “Swiss Army infinitesimal jackknife,” like the famous Swiss Army knife, is a single tool with many different functions.

2 Methods and Results

2.1 Problem definition

We consider the problem of estimating an unknown parameter $\theta \in \Omega_\theta \subseteq \mathbb{R}^D$, with a compact Ω_θ and a dataset of size N . Our analysis will proceed entirely in terms of a fixed dataset, though we will be careful to make assumptions that will plausibly hold for all N under suitably well-behaved random sampling. We define our estimate, $\hat{\theta} \in \Omega_\theta$, as the root of a weighted estimating equation. For each $n = 1, \dots, N$, let $g_n(\theta)$ be a function from Ω_θ to \mathbb{R}^D . Let w_n be a real number, and let w be the vector collecting the w_n . Then $\hat{\theta}$ is defined as the quantity that satisfies

$$\hat{\theta}(w) := \theta \text{ such that } \frac{1}{N} \sum_{n=1}^N w_n g_n(\theta) = 0. \quad (1)$$

We will impose assumptions below that imply at least local uniqueness of $\hat{\theta}(w)$; see the discussion following Assumption 2 in Section 2.3.

As an example, consider a family of continuously differentiable loss functions $f(\cdot, \theta)$ parameterized by θ and evaluated at data points $x_n, n = 1, \dots, N$. If we want to solve the optimization problem $\hat{\theta} = \operatorname{argmin}_{\theta \in \Omega_\theta} \frac{1}{N} \sum_{n=1}^N f(x_n, \theta)$, then we take $g_n(\theta) = \partial f(x_n, \theta) / \partial \theta$ and $w_n \equiv 1$. By keeping our notation general, we will be able to analyze a more general class of problems, such as multi-stage optimization (see Section 6). However, to aid intuition, we will sometimes refer to the $g_n(\theta)$ as “gradients” and their derivatives as “Hessians.”

When equation (1) is not degenerate (we articulate precise conditions below), $\hat{\theta}$ is a function of the weights through solving the estimating equation, and we

write $\hat{\theta}(w)$ to emphasize this. We will focus on the case where we have solved equation (1) for the weight vector of all ones, $1_w := (1, \dots, 1)$, which we denote $\hat{\theta}_1 := \hat{\theta}(1_w)$.

A re-sampling scheme can be specified by choosing a set $W \subseteq \mathbb{R}^N$ of weight vectors. For example, to approximate leave- k -out CV, one repeatedly computes $\hat{\theta}(w)$ where w has k randomly chosen zeros and all ones otherwise. Define W_k as the set of every possible leave- k -out weight vector. Showing that our approximation is good for all leave- k -out analyses with probability one is equivalent to showing that the approximation is good for all $w \in W_k$.

In the case of the bootstrap, W contains a fixed number B of randomly chosen weight vectors, $w_b^* \stackrel{iid}{\sim} \text{Multinomial}(N, N^{-1})$ for $b = 1, \dots, B$, so that $\sum_{n=1}^N w_{bn}^* = N$ for each b . Note that while w_n or w_{bn}^* are scalars, w_b^* is a vector of length N . The distribution of $\hat{\theta}(w_b^*) - \hat{\theta}(1_w)$ is then used to estimate the sampling variation of $\hat{\theta}_1$. Define this set $W_B^* = \{w_1^*, \dots, w_B^*\}$. Note that W_B^* is stochastic and is a subset of all weight vectors that sum to N .

In general, W can be deterministic or stochastic, may contain integer or non-integer values, and may be determined independently of the data or jointly with it. As with the data, our results hold for a given W , but in a way that will allow natural high-probability extensions to stochastic W .

2.2 Linear approximation

The main problem we solve is the computational expense involved in evaluating $\hat{\theta}(w)$ for all the $w \in W$. Our contribution is to use only quantities calculated from $\hat{\theta}_1$ to approximate $\hat{\theta}(w)$ for all $w \in W$, without re-solving equation (1). Our approximation is based on the derivative $\frac{d\hat{\theta}(w)}{dw^T}$, whose existence depends on the derivatives of $g_n(\theta)$, which we assume to exist, and which we denote as $h_n(\theta) := \frac{\partial g_n(\theta)}{\partial \theta^T}$. We use this notation because $h_n(\theta)$ would be the Hessian of a term of the objective in the case of an optimization problem. We make the following definition for brevity.

Definition 1. The fixed point equation and its derivative are given respectively by

$$G(\theta, w) := \frac{1}{N} \sum_{n=1}^N w_n g_n(\theta)$$

$$H(\theta, w) := \frac{1}{N} \sum_{n=1}^N w_n h_n(\theta).$$

Note that $G(\hat{\theta}(w), w) = 0$ because $\hat{\theta}(w)$ solves equation (1) for w . We define $H_1 := H(\hat{\theta}_1, 1_w)$ and define the weight difference as $\Delta w = w - 1_w \in \mathbb{R}^N$. When H_1 is invertible, one can use the implicit function theorem and the chain rule to

show that the derivative of $\hat{\theta}(w)$ with respect to w is given by

$$\begin{aligned}\frac{d\hat{\theta}(w)}{dw^T}|_{1_w}\Delta w &= -H_1^{-1}\frac{1}{N}\sum_{n=1}^N g_n(\hat{\theta}_1)\Delta w \\ &= -H_1^{-1}G(\hat{\theta}_1, \Delta w).\end{aligned}$$

This derivative allows us to form a first-order approximation to $\hat{\theta}(w)$ at $\hat{\theta}_1$.

Definition 2. Our linear approximation to $\hat{\theta}(w)$ is given by

$$\hat{\theta}_{\text{IJ}}(w) := \hat{\theta}_1 - H_1^{-1}G(\hat{\theta}_1, \Delta w).$$

We use the subscript “IJ” for “infinitesimal jackknife,” which is the name for this estimate in the statistics literature [Jaekel, 1972, Shao, 1993]. Because $\hat{\theta}_{\text{IJ}}$ depends only on $\hat{\theta}_1$ and Δw , and not on solutions at any other values of w , there is no need to re-solve equation (1). Instead, to calculate $\hat{\theta}_{\text{IJ}}$ one must solve a linear system involving H_1 . Recalling that θ is D -dimensional, the calculation of H_1^{-1} (or a factorization that supports efficient solution of linear systems) can be $O(D^3)$. However, once H_1^{-1} is calculated or H_1 is factorized, calculating our approximation $\hat{\theta}_{\text{IJ}}(w)$ for each new weight costs only as much as a single matrix-vector multiplication. Furthermore, H_1 often has a sparse structure allowing H_1^{-1} to be calculated more efficiently than a worst-case scenario (see Section 6 for an example). In more high-dimensional examples with dense Hessian matrices, such as neural networks, one may need to turn to approximations such as stochastic second-order methods [Koh and Liang, 2017, Agarwal et al., 2017] and conjugate gradient [Wright and Nocedal, 1999]. Indeed, even in relatively small or sparse problems, the vast bulk of the computation required to calculate $\hat{\theta}_{\text{IJ}}$ is in the computation of H_1^{-1} . We leave the important question of approximate calculation of H_1^{-1} for future work.

2.3 Assumptions and results

We now state our key assumptions and results, which are sufficient conditions under which $\hat{\theta}_{\text{IJ}}(w)$ will be a good approximation to $\hat{\theta}(w)$. We defer most proofs to Appendix A. We use $\|\cdot\|_{op}$ to denote the matrix operator norm, $\|\cdot\|_2$ to denote the L_2 norm, and $\|\cdot\|_1$ to denote the L_1 norm. For quantities like g and h , which have dimensions $N \times D$ and $N \times D \times D$ respectively, we apply the L_p norm to the vectorized version of arrays. For example, $\frac{1}{\sqrt{N}}\|h(\theta)\|_2 = \sqrt{\frac{1}{N}\sum_{n=1}^N\sum_{i=1}^D\sum_{j=1}^D[h_n(\theta)]_{ij}^2}$ which is the square root of a sample average over $n \in [N]$.

We state all assumptions and results for a fixed N , a given estimating equation vector $g(\theta)$, and a fixed class of weights W . Although our analysis proceeds with these quantities fixed, we are careful to make only assumptions that can plausibly hold for all N and/or for randomly chosen W under appropriate regularity conditions.

Assumption 1 (Smoothness). *For all $\theta \in \Omega_\theta$, each $g_n(\theta)$ is continuously differentiable in θ .*

The smoothness in Assumption 1 is necessary for a local approximation like Definition 2 to have any hope of being useful.

Assumption 2 (Non-degeneracy). *For all $\theta \in \Omega_\theta$, $H(\theta, 1_w)$ is non-singular, with $\sup_{\theta \in \Omega_\theta} \|H(\theta, 1_w)^{-1}\|_{op} \leq C_{op} < \infty$.*

Without Assumption 2, the derivative in Definition 2 would not exist. For an optimization problem, Definition 2 amounts to assuming that the Hessian is strongly positive definite, and, in general, assures that the solution $\hat{\theta}_1$ is unique. Under our assumptions, we will show later that, additionally, $\hat{\theta}(w)$ is unique in a neighborhood of $\hat{\theta}_1$; see Lemma 6 of Appendix A. Furthermore, by fixing C_{op} , if we want to apply Assumption 2 for $N \rightarrow \infty$, we will require that H_1 remains strongly positive definite.

Assumption 3 (Bounded averages). *There exist finite constants C_g and C_h such that $\sup_{\theta \in \Omega_\theta} \frac{1}{\sqrt{N}} \|g(\theta)\|_2 \leq C_g < \infty$ and $\sup_{\theta \in \Omega_\theta} \frac{1}{\sqrt{N}} \|h(\theta)\|_2 \leq C_h < \infty$.*

Assumption 3 essentially states that the sample variances of the gradients and Hessians are uniformly bounded. Note that it does not require that these quantities are bounded term-wise. For example, we allow $\sup_n \|g_n(\theta)\|_2^2 \xrightarrow{N \rightarrow \infty} \infty$, as long as $\sup_n \frac{1}{N} \|g_n(\theta)\|_2^2$ remains bounded. This is a key advantage of the present work over many past applications of the IJ to M-estimation, which require $\sup_n \|g_n(\theta)\|_2^2$ to be uniformly bounded for all N [Shao and Tu, 2012, Beirami et al., 2017].

In both machine learning and statistics, $\sup_n \|g_n(\theta)\|_2^2$ is rarely bounded, though $\frac{1}{N} \|g(\theta)\|_2^2$ often is. As a simple example, suppose that $\theta \in \mathbb{R}^1$, $x_n \sim \mathcal{N}(0, 1)$, and $g_n = \theta - x_n$, as would arise from the squared error loss $f_n(x_n, \theta) = \frac{1}{2}(\theta - x_n)^2$. Fix a θ and let $N \rightarrow \infty$. Then $\sup_n \|g_n(\theta)\|_2^2 \rightarrow \infty$ because $\sup_n |x_n| \rightarrow \infty$, but $\frac{1}{N} \|g(\theta)\|_2^2 \rightarrow \theta^2 + 1$ by the law of large numbers.

Assumption 4 (Local smoothness). *There exists a $\Delta_\theta > 0$ and a finite constant L_h such that, $\|\theta - \hat{\theta}_1\|_2 \leq \Delta_\theta$ implies that $\frac{\|h(\theta) - h(\hat{\theta}_1)\|_2}{\sqrt{N}} \leq L_h \|\theta - \hat{\theta}_1\|_2$.*

The constants defined in Assumption 4 are needed to calculate our error bounds explicitly.

Assumptions 1–4 are quite general and should be expected to hold for many reasonable problems, including holding uniformly asymptotically with high probability for many reasonable data-generating distributions, as the following lemma shows.

Lemma 1 (The assumptions hold under uniform convergence). *Let Ω_θ be a compact set, and let $g_n(\theta)$ be twice continuously differentiable IID random functions for $n \in [N]$. (The function is random but θ is not—for example,*

$\mathbb{E}[g_n(\theta)]$ is still a function of θ .) Define $r_n(\theta) := \frac{\partial^2 g_n(\theta)}{\partial \theta \partial \theta}$, so $r_n(\theta)$ is a $D \times D \times D$ tensor.

Assume that we can exchange integration and differentiation, that $\mathbb{E}[h_n(\theta)]$ is non-singular for all $\theta \in \Omega_\theta$, and that all of $\mathbb{E}\left[\sup_{\theta \in \Omega_\theta} \|g_n(\theta)\|_2^2\right]$, $\mathbb{E}\left[\sup_{\theta \in \Omega_\theta} \|h_n(\theta)\|_2^2\right]$, and $\mathbb{E}\left[\sup_{\theta \in \Omega_\theta} \|r_n(\theta)\|_2^2\right]$ are finite.

Then $\lim_{N \rightarrow \infty} P(\text{Assumptions 1-4 hold}) = 1$.

Lemma 1 follows from the uniform convergence results of Theorems 9.1 and 9.2 in Keener [2011]. See Appendix A.4 for a detailed proof. A common example to which Lemma 1 would apply is where x_n are well-behaved IID data and $g_n(\theta) = \gamma(x_n, \theta)$ for an appropriately smooth estimating function $\gamma(\cdot, \theta)$. See Keener [2011, Chapter 9] for more details and examples, including applications to maximum likelihood estimators on unbounded domains.

Assumptions 1–4 apply to the estimating equation. We also require a boundedness condition for W .

Assumption 5 (Bounded weight averages). *The quantity $\frac{1}{\sqrt{N}} \|w\|_2$ is uniformly bounded for $w \in W$ by a finite constant C_w .*

Our final requirement is considerably more restrictive, and contains the essence of whether or not $\hat{\theta}_{\text{IJ}}(w)$ will be a good approximation to $\hat{\theta}(w)$.

Condition 1 (Set complexity). *There exists a $\delta \geq 0$ and a corresponding set $W_\delta \subseteq W$ such that*

$$\begin{aligned} \sup_{w \in W_\delta} \sup_{\theta \in \Omega_\theta} \left\| \frac{1}{N} \sum_{n=1}^N (w_n - 1) g_n(\theta) \right\|_1 &\leq \delta \quad \text{and} \\ \sup_{w \in W_\delta} \sup_{\theta \in \Omega_\theta} \left\| \frac{1}{N} \sum_{n=1}^N (w_n - 1) h_n(\theta) \right\|_1 &\leq \delta. \end{aligned}$$

Condition 1 is central to establishing when the approximation $\hat{\theta}_{\text{IJ}}(w)$ is accurate. For a given δ , W_δ will be the class of weight vectors for which $\hat{\theta}_{\text{IJ}}(w)$ is accurate to within order δ . Trivially, $1_w \in W_\delta$ for $\delta = 0$, so W_δ is always non-empty, even for arbitrarily small δ . The trick will be to choose a small δ that still admits a large class W_δ of weight vectors. In Section 3 we will discuss Condition 1 in more depth, but it will help to first state our main theorem.

Definition 3. The following constants are given by quantities in Assumptions 1–5.

$$\begin{aligned} C_{\text{IJ}} &:= 1 + DC_w L_h C_{op} \\ \Delta_\delta &:= \min \left\{ \Delta_\theta C_{op}^{-1}, \frac{1}{2} C_{\text{IJ}}^{-1} C_{op}^{-1} \right\}. \end{aligned}$$

Note that, although the parameter dimension D occurs explicitly only once in Definition 3, all of C_w , C_{op} , and L_h in general might also contain dimension

dependence. Additionally, the bound δ in Condition 1, a measure of the set complexity of the parameters, will typically depend on dimension. However, the particular place where the parameter dimension enters will depend on the problem and asymptotic regime, and our goal is to provide an adaptable toolkit for a wide variety of problems.

We are now ready to state our main result.

Theorem 1 (Error bound for the approximation). *Under Assumptions 1–5 and Condition 1,*

$$\delta \leq \Delta_\delta \Rightarrow \max_{w \in W_\delta} \left\| \hat{\theta}_{IJ}(w) - \hat{\theta}(w) \right\|_2 \leq 2C_{op}^2 C_{IJ} \delta^2.$$

We stress that Theorem 1 bounds only the difference between $\hat{\theta}_{IJ}(w)$ and $\hat{\theta}(w)$. Theorem 1 alone does not guarantee that $\hat{\theta}_{IJ}(w)$ converges to any hypothetical infinite population quantity. We see this as a strength, not a weakness. To begin with, convergence to an infinite population requires stronger assumptions. Contrast, for example, the Fréchet differentiability work of Clarke [1983], on which our work is based, with the stricter requirements in the proof of consistency in Shao [1993]. Second, machine learning problems may not naturally admit a well-defined infinite population, and the dataset at hand may be of primary interest. Finally, by analyzing a particular sample rather than a hypothetical infinite population, we can bound the error in terms of the quantities C_{IJ} and Δ_δ , which can actually be calculated from the data at hand.

Still, Theorem 1 is useful to prove asymptotic results about the difference $\left\| \hat{\theta}_{IJ}(w) - \hat{\theta}(w) \right\|_2$. As an illustration, we now show that the uniform consistency of leave- k -out CV follows from Theorem 1 by a straightforward application of Hölder’s inequality.

Corollary 1 (Consistency for leave- k -out CV). *Assume that Assumptions 1–5 hold uniformly for all N . Fix an integer k , and let*

$$W_k := \{w : w_n = 0 \text{ in } k \text{ entries and } 1 \text{ otherwise}\}.$$

Then, for all N , there exists a constant C_K such that

$$\begin{aligned} \sup_{w \in W_k} \left\| \hat{\theta}_{IJ}(w) - \hat{\theta}(w) \right\|_2 &\leq C_K \frac{\|g\|_\infty^2}{N^2} \\ &\leq C_K \frac{\max\{C_g, C_h\}^2}{N}. \end{aligned}$$

Proof. For $w \in W_k$, $\frac{\|\Delta w\|_2}{\sqrt{N}} = \sqrt{\frac{k}{N}}$. Define $C_{gh} := \max\{C_g, C_h\}$. By Assumption 3, $\|g\|_2 / \sqrt{N} \leq C_{gh}$ and $\|h\|_2 / \sqrt{N} \leq C_{gh}$ for all N . By Hölder’s inequality,

$$\begin{aligned} &\sup_{w \in W} \sup_{\theta \in \Omega_\theta} \left\| \frac{1}{N} \sum_{n=1}^N (w_n - 1) g_n(\theta) \right\|_1 \\ &\leq \sup_{w \in W} \|w - 1_w\|_1 \sup_{\theta \in \Omega_\theta} \frac{\|g\|_\infty}{N} = K \frac{\|g\|_\infty}{N} \leq K \frac{C_{gh}}{\sqrt{N}}, \end{aligned}$$

with a similar bound for $\|h\|_2$. Consequently, for N large enough, Condition 1 is satisfied with $W_\delta = W_k$ and either $\delta = K \frac{\|g\|_\infty}{N}$ or $\delta = K \frac{C_{gh}}{\sqrt{N}}$. The result then follows from Theorem 1. \square

3 Examples

The moral of Theorem 1 is that, under Assumptions 1–5 and Condition 1, $\|\hat{\theta}_{IJ} - \hat{\theta}(w)\| = O(\delta^2)$ for $w \in W_\delta$. That is, if we can make δ small enough, W_δ big enough, and still satisfy Condition 1, then $\hat{\theta}_{IJ}(w)$ is a good approximation to $\hat{\theta}(w)$ for “most” w , where “most” is defined as the size of W_δ . So it is worth taking a moment to develop some intuition for Condition 1. We have already seen in Corollary 1 that $\hat{\theta}_{IJ}$ is, asymptotically, a good approximation for leave- k -out CV uniformly in W . We now discuss some additional cases: first, a worst-case example for which $\hat{\theta}_{IJ}$ is not expected to work, second the bootstrap, and finally we revisit leave-one-out cross validation in the context of these other two methods.

First, consider a pathological example. Let W_{full} be the set of all weight vectors that sum to N . Let $n^* = \max_{n \in [N]} \|g_n(\hat{\theta}_1)\|_1$ be the index of the gradient term with the largest L_1 norm, and let $w_{n^*} = N$ and $w_n = 0$ for $n \neq n^*$. Then

$$\begin{aligned} & \sup_{\theta \in \Omega_\theta} \left\| \frac{1}{N} \sum_{n=1}^N (w_n - 1) g_n(\theta) \right\|_1 \\ &= \sup_{\theta \in \Omega_\theta} \left\| g_{n^*}(\theta) - \frac{1}{N} \sum_{n=1}^N g_n(\theta) \right\|_1 \geq \|g_{n^*}(\hat{\theta}_1)\|_1. \end{aligned}$$

(The last inequality uses the fact that $G(\hat{\theta}_1, 1_w) = 0$.) In this case, unless the largest gradient, $\|g_{n^*}(\hat{\theta}_1)\|_1$, is small, Condition 1 will not be satisfied for small δ , and we would not expect $\hat{\theta}_{IJ}$ to be a good estimate for $\hat{\theta}(w)$ for all $w \in W_{full}$. The class W_{full} is too expressive. In the language of Condition 1, for some small fixed δ , W_δ will be some very restricted subset of W_{full} in most realistic situations.

Now, suppose that we are using B bootstrap weights, $w_b^* \stackrel{iid}{\sim} \text{Multinomial}(N, N^{-1})$ for $b = 1, \dots, B$, and analyzing an optimization problem as defined in Section 2.1. For a given w_b^* , a dataset x_1^*, \dots, x_N^* formed by taking $w_{b,n}^*$ copies of datapoint x_n is equivalent in distribution to N IID samples with replacement from the

empirical distribution on (x_1, \dots, x_N) . In this notation, we then have

$$\begin{aligned} \frac{1}{N} \sum_{n=1}^N (w_b^* - 1) g_n(\theta) = \\ \frac{1}{N} \sum_{n=1}^N \frac{\partial f(\theta, x_n^*)}{\partial \theta} - \frac{1}{N} \sum_{n=1}^N \frac{\partial f(\theta, x_n)}{\partial \theta}. \end{aligned}$$

In this case, Condition 1 is a uniform bound on a centered empirical process of derivatives of the objective function. Note that estimating sample variances by applying the IJ with bootstrap weights is equivalent to the ordinary delta method based on an asymptotic normal approximation [Efron, 1982, Chapter 21]. In order to provide an approximation to the bootstrap that retains benefits (such as the faster-than-normal convergence to the true sampling distribution described by Hall [2013]), one must consider higher-ordered Taylor expansions of $\hat{\theta}(w)$. We leave this for future work.

Finally, let us return to leave-one-out CV. In this case, $w_n - 1$ is nonzero for exactly one entry. Again, we can choose to leave out the adversarially-chosen n^* as in the first pathological example. However, unlike the pathological example, the leave-one-out CV weights are constrained to be closer to 1_w —specifically, we set $w_{n^*} = 0$, and let w be one elsewhere. Then Condition 1 requires $\sup_{\theta \in \Omega_\theta} \left\| \frac{1}{N} g_{n^*}(\theta) \right\|_1 \leq \delta$. In contrast to the pathological example, this supremum will get smaller as N increases as long as $\|g_{n^*}(\theta)\|_1$ grows more slowly than N . For this reason, we expect leave-one-out (and, indeed, leave- k -out for fixed k) to be accurately approximated by $\hat{\theta}_{\text{IJ}}$ in many cases of interest, as stated in Corollary 1.

4 Related Work

Although the idea of forming a linear approximation to the re-weighting of an M-estimator has a long history, we nevertheless contribute in a number of ways. By limiting ourselves to approximating the exact reweighting on a particular dataset, we both loosen the strict requirements from the statistical literature and generalize the existing results from the machine learning literature.

The jackknife is often favored over the IJ in the statistics literature because of the former’s simple computational approach, as well as perceived difficulties in calculating the necessary derivatives when some of the parameters are implicitly defined via optimization [Shao and Tu, 2012, Chapter 2.1] (though exceptions exist; see, e.g., Wager et al. [2014]). The brute-force approach of the jackknife is, however, a liability in large-scale machine learning problems, which are generally extremely expensive to re-optimize. Furthermore, and critically, the complexity and tedium of calculating the necessary derivatives is entirely eliminated by modern automatic differentiation [Baydin et al., 2017, Maclaurin et al., 2015].

Our work is based on the proof of the Fréchet differentiability of M-estimators of Clarke [1983]. In classical statistics, Fréchet differentiability is typically used

to describe the asymptotic behavior of functionals of the empirical distribution in terms of a functional [Mises, 1947, Fernholz, 1983]. Since Clarke [1983] was motivated by such asymptotic questions, he studied the Fréchet derivative evaluated at a continuous probability distribution for function classes that included delta functions. This focus led to the requirement of a bounded gradient. However, unbounded gradients are ubiquitous in both statistics and machine learning, and an essential contribution of the current paper is to remove the need for bounded gradients.

There exist proofs of the consistency of the (non-infinitesimal) jackknife that allow for unbounded gradients. For example, it is possible that the proofs of Reeds [1978], which require a smoothness assumption similar to our Assumption 4, could be adapted to the IJ. However, the results of Reeds [1978]—as well as those of Clarke [1983] and subsequent applications such as those of Shao and Tu [2012]—are asymptotic and applicable only to IID data. By providing finite sample results for a fixed dataset and weight set, we are able to provide a template for proving accuracy bounds for more generic probability distributions and re-weighting schemes.

A number of recent machine learning papers have derived approximate linear versions of leave-one-out estimators. Koh and Liang [2017] consider approximating the effect of leaving out one observation at a time to discover influential observations and construct adversarial examples, but provide little supporting theory. Beirami et al. [2017] provide rigorous proofs for an approximate leave-one-out CV estimator; however, their estimator requires computing a new inverse Hessian for each new weight at the cost of a considerable increase in computational complexity. Like the classical statistics literature, Beirami et al. [2017] assume that the gradients are bounded for all N . When $\|g\|_\infty^2$ in Corollary 1 is finite for all N , we achieve the same N^{-2} rate claimed by Beirami et al. [2017] for leave-one-out CV although we use only a single matrix inverse. Rad and Maleki [2018] also approximate leave-one-out CV, and prove tighter bounds for the error of their approximation than we do, but their work is customized to leave-one-out CV and makes much more restrictive assumptions (e.g., Gaussianity).

5 Simulated Experiments

We begin the empirical demonstration of our method on two simple generalized linear models: logistic and Poisson regression.¹ In each case, we generate a synthetic dataset $Z = \{(x_n, y_n)\}_{n=1}^N$ from parameters (θ, b) , where $\theta \in \mathbb{R}^{100}$ is a vector of regression coefficients and $b \in \mathbb{R}$ is a bias term. In each experiment, $x_n \in \mathbb{R}^{100}$ is drawn from a multivariate Gaussian, and y_n is a scalar drawn from a Bernoulli distribution with the logit link or from a Poisson distribution with the exponential link.

¹Leave-one-out CV may not be the most appropriate estimator of generalization error in this setting [Rosset and Tibshirani, 2018], but this section is intended only to provide simple illustrative examples.

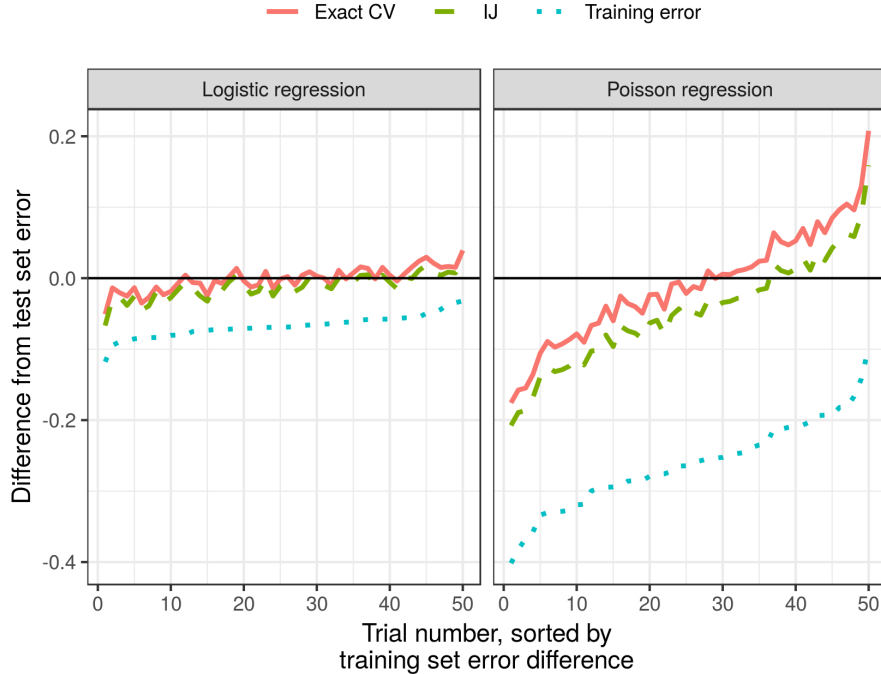


Figure 1: Simulated data: accuracy results.

For a ground truth, we generate a large test set with $N = 100,000$ datapoints to measure the true generalization error. We show in Fig. 1 that, over 50 randomly generated datasets, our approximation consistently underestimates the actual error predicted by exact leave-one-out CV; however, the difference is small relative to the improvements they both make over the error evaluated on the training set.

Fig. 2 shows the relative timings of our approximation and exact leave-one-out CV on logistic regression with datasets of increasing size. The time to run our approximation is roughly an order of magnitude smaller.

6 Genomics Experiments

We now consider a genomics application in which we use CV to choose the degree of a spline smoother when clustering time series of gene expression data. Code and instructions to reproduce our results can be found in the git repository [rjiordan/AISTATS2019SwissArmyIJ](https://github.com/rjiordan/AISTATS2019SwissArmyIJ). The application is also described in detail in Appendix B.

We use a publicly available data set of mice gene expression [Shoemaker et al., 2015] in which mice were infected with influenza virus, and gene expression was




Figure 2: Simulated data: timing results.

assessed several times after infection. The observed data consists of expression levels y_{gt} for genes $g = 1, \dots, n_g$ and time points $t = 1, \dots, n_t$. In our case $n_g = 1000$ and $n_t = 14$. Many genes behave the same way; thus, clustering the genes by the pattern of their behavior over time allows dimensionality reduction that can facilitate interpretation. Consequently, we wish to first fit a smoothed regression line to each gene and then cluster the results. Following Luan and Li [2003], we model the time series as a gene-specific constant additive offset plus a B-spline basis of degree 3, and the task is to choose the B-spline basis degrees of freedom using cross-validation on the time points.

Our analysis runs in two stages—first, we regress the genes on the spline basis, and then we cluster a transformed version of the regression fits. By modeling in two stages, we both speed up the clustering and allow for the use of flexible transforms of the fits. We are interested in choosing the smoothing parameter using CV on the time points. Both the time points and the smoothing parameter enter the regression objective directly, but they affect the clustering objective only through the optimal regression parameters. Because the optimization proceeds in two stages, the fit is not the optimum of any single objective function. However, it can still be represented as an M-estimator (see Appendix B).

We implemented the model in `scipy` [Jones et al., 2001] and computed all derivatives with `autograd` [Maclaurin et al., 2015]. We note that the match between “exact” cross-validation (removing time points and re-optimizing) and the IJ was considerably improved by using a high-quality second-order optimization method. In particular, for these experiments, we employed the Newton conjugate-gradient trust region method [Wright and Nocedal, 1999, Chapter 7.1] as implemented by the method `trust-ncg` in `scipy.optimize`, preconditioned by the Cholesky decomposition of an inverse Hessian calculated at an initial

approximate optimum. The Hessian used for the preconditioner was with respect to the clustering parameters only and so could be calculated quickly, in contrast to the H_1 matrix used for the IJ, which includes the regression parameters as well. We found that first-order or quasi-Newton methods (such as BFGS) often got stuck or terminated at points with fairly large gradients. At such points our method does not apply in theory nor, we found, very well in practice.




Figure 3: Genomics data: accuracy results.

Fig. 3 shows that the IJ is a reasonably good approximation to the test set error.² In particular, both the IJ and exact CV capture the increase in test error for $df = 8$, which is not present in the training error. Thus we see that, like exact CV, the IJ is able to prevent overfitting. Though the IJ underestimates exact CV, we note that it differs from exact CV by no more than exact CV itself differs from the true quantity of interest, the test error.

The timing results for the genomics experiment are shown in Fig. 4. For this particular problem with approximately 39,000 parameters (the precise number depends on the degrees of freedom), finding the initial optimum takes about 42 seconds. The cost of finding the initial optimum is shared by exact CV and the

²In fact, in this case, the IJ is a better predictor of test set error than exact CV. However, the authors have no reason at present to believe that the IJ is a better predictor of test error than exact CV in general.




Figure 4: Genomics data: timing results.

IJ, and, as shown in Fig. 4, is a small proportion of both.

The principle time cost of the IJ is the computation of H_1 . Computing and inverting a dense matrix of size 39,000 would be computationally prohibitive. But, for the regression objective, H_1 is extremely sparse and block diagonal, so computing H_1 in this case took only around 360 seconds. Inverting H_1 took negligible time. Once we have H_1^{-1} , obtaining the subsequent IJ approximations is nearly instantaneous.

The cost of refitting the model for exact CV varies by degrees of freedom (increasing degrees of freedom increases the number of parameters) and the number of left-out points (an increasing number of left-out datapoints increases the number of refits). As can be seen in Fig. 4, for low degrees of freedom and few left-out points, the cost of re-optimizing is approximately the same as the cost of computing H_1 . However, as the degrees of freedom and number of left-out points grow, the cost of exact CV increases to as much as an order of magnitude more than that of the IJ.

7 Conclusion

We recommend consideration of the Swiss Army infinitesimal jackknife for modern machine learning problems. The large size of modern data both increases the need for fast approximations and renders such approximations more accurate.

Furthermore, modern automatic differentiation renders many past practical difficulties obsolete. By stepping back from the strict requirements of classical statistical theory, we can see that the value of the infinitesimal jackknife extends beyond its traditional application areas, while retaining desirable generality in other respects.

Acknowledgements. We thank anonymous reviewers for their helpful comments and suggestions. We are grateful to Nelle Varoquaux for her help with the genomics experiments. This research was supported in part by DARPA (FA8650-18-2-7832), an ARO YIP award, an NSF CAREER award, and the CSAIL-MSR Trustworthy AI Initiative. Ryan Giordano was supported by the Gordon and Betty Moore Foundation through Grant GBMF3834 and by the Alfred P. Sloan Foundation through Grant 2013-10-27 to the University of California, Berkeley. Runjing Liu was supported by the NSF Graduate Research Fellowship.

References

- N. Agarwal, B. Bullins, and E. Hazan. Second-order stochastic optimization in linear time. *Journal of Machine Learning Research*, 2017.
- A. Baydin, B. Pearlmutter, A. Radul, and J. Siskind. Automatic differentiation in machine learning: a survey. *Journal of Machine Learning Research*, 18(153): 1–153, 2017.
- A. Beirami, M. Razaviyayn, S. Shahrampour, and V. Tarokh. On optimal generalizability in parametric learning. In *Advances in Neural Information Processing Systems (NIPS)*, pages 3458–3468, 2017.
- B. Clarke. Uniqueness and Fréchet differentiability of functional solutions to maximum likelihood type equations. *The Annals of Statistics*, 11(4):1196–1205, 1983.
- R. Dudley. *Real analysis and probability*. Chapman and Hall/CRC, 2018.
- B. Efron. *The Jackknife, the Bootstrap, and Other Resampling Plans*, volume 38. Society for Industrial and Applied Mathematics, 1982.
- L. Fernholz. *Von Mises Calculus for Statistical Functionals*, volume 19. Springer Science & Business Media, 1983.
- P. Hall. *The Bootstrap and Edgeworth Expansion*. Springer Science & Business Media, 2013.
- L. Jaeckel. The infinitesimal jackknife, memorandum. Technical report, MM 72-1215-11, Bell Lab. Murray Hill, NJ, 1972.
- E. Jones, T. Oliphant, P. Peterson, et al. SciPy: Open source scientific tools for Python, 2001. URL <http://www.scipy.org/>.
- R. W. Keener. *Theoretical Statistics: Topics for a Core Course*. Springer, 2011.
- P. W. Koh and P. Liang. Understanding black-box predictions via influence functions. In *International Conference on Machine Learning (ICML)*, 2017.
- Y. Luan and H. Li. Clustering of time-course gene expression data using a mixed-effects model with B-splines. *Bioinformatics*, 19(4):474–482, 2003.

- D. Maclaurin, D. Duvenaud, and R. P. Adams. Autograd: Effortless gradients in numpy. In *International Conference on Machine Learning (ICML) AutoML Workshop*, 2015.
- R. Mises. On the asymptotic distribution of differentiable statistical functions. *The Annals of Mathematical Statistics*, 18(3):309–348, 1947.
- K. R. Rad and A. Maleki. A scalable estimate of the extra-sample prediction error via approximate leave-one-out. *arXiv Preprint*, January 2018.
- J. A. Reeds. Jackknifing maximum likelihood estimates. *The Annals of Statistics*, pages 727–739, 1978.
- S. Rosset and R. J. Tibshirani. From fixed-X to random-X regression: Bias-variance decompositions, covariance penalties, and prediction error estimation. *Journal of the American Statistical Association*, 2018.
- J. Schott. *Matrix Analysis for Statistics*. John Wiley & Sons, 2016.
- J. Shao. Differentiability of statistical functionals and consistency of the jackknife. *The Annals of Statistics*, pages 61–75, 1993.
- J. Shao and D. Tu. *The Jackknife and Bootstrap*. Springer Series in Statistics, 2012.
- J. E. Shoemaker, S. Fukuyama, A. J. Eisfeld, D. Zhao, E. Kawakami, S. Sakabe, T. Maemura, T. Gorai, H. Katsura, Y. Muramoto, S. Watanabe, T. Watanabe, K. Fuji, Y. Matsuoka, H. Kitano, and Y. Kawaoka. An ultrasensitive mechanism regulates influenza virus-induced inflammation. *PLoS Pathogens*, 11(6):1–25, 2015.
- S. Wager, T. Hastie, and B. Efron. Confidence intervals for random forests: The jackknife and the infinitesimal jackknife. *The Journal of Machine Learning Research*, 15(1):1625–1651, 2014.
- S. Wright and J. Nocedal. *Numerical Optimization*, volume 35. 1999.

A Detailed assumptions, lemmas, and proofs

A.1 Tools

We begin by stating two general propositions that will be useful. First, we show that a version of Cauchy-Schwartz can be applied to weighted sums of tensors.

Proposition 1. *Tensor array version of Hölder's inequality. Let w be an array of scalars and let $a = (a_1, \dots, a_N)$ be an array of tensors, where each a_n is indexed by $i = 1, \dots, D_A$ (i may be a multi-index—e.g., if A is a $D \times D$ matrix, then $i = (j, k)$, for $j, k \in [D]$ and $D_A = D^2$). Let $p, q \in [1, \infty]$ be two numbers such that $p^{-1} + q^{-1} = 1$. Then*

$$\left\| \frac{1}{N} \sum_{n=1}^N w_n a_n \right\|_1 \leq \frac{D_A^{\frac{1}{p}}}{N} \|w\|_p \|a\|_q.$$

In particular, with $p = q = 2$,

$$\left\| \frac{1}{N} \sum_{n=1}^N w_n a_n \right\|_1 \leq \sqrt{D_A} \frac{\|w\|_2}{\sqrt{N}} \frac{\|a\|_2}{\sqrt{N}}.$$

Proof. The conclusion follows from the ordinary Hölder's inequality applied term-wise to n and Jensen's inequality applied to the indices i .

$$\begin{aligned} \left\| \frac{1}{N} \sum_{n=1}^N w_n a_n \right\|_1 &= \sum_{i=1}^{D_A} \left| \frac{1}{N} \sum_{n=1}^N w_n (a_n)_i \right| \\ &\leq \frac{1}{N} \sum_{i=1}^{D_A} \left| \left(\sum_{n=1}^N |w_n|^p \right)^{\frac{1}{p}} \left(\sum_{n=1}^N |(a_n)_i|^q \right)^{\frac{1}{q}} \right| \quad (\text{Hölder}) \\ &= \frac{1}{N} \|w\|_p \frac{D_A}{D_A} \sum_{i=1}^{D_A} \left(\sum_{n=1}^N |(a_n)_i|^q \right)^{\frac{1}{q}} \\ &\leq \frac{1}{N} \|w\|_p D_A \left(\frac{1}{D_A} \sum_{i=1}^{D_A} \sum_{n=1}^N |(a_n)_i|^q \right)^{\frac{1}{q}} \quad (\text{Jensen applied to } i) \\ &= \frac{1}{N} \|w\|_p D_A \left(\frac{1}{D_A} \sum_{n=1}^N \|a_n\|_q^q \right)^{\frac{1}{q}} \\ &= \frac{1}{N} \|w\|_p D_A^{1-\frac{1}{q}} \|a\|_q \\ &= \frac{D_A^{\frac{1}{p}}}{N} \|w\|_p \|a\|_q. \end{aligned}$$

□

Next, we prove a relationship between the term-wise difference between matrices and the difference between their operator norms. It is well-known that the minimum eigenvalue of a non-singular matrix is continuous in the entries of the matrix. In the next proposition, we quantify this continuity for the L_1 norm.

Proposition 2. *Let A and B be two square matrices of the same size. Let $\|A^{-1}\|_{op} \leq C_{op}$ for some finite C_{op} , and Then*

$$\|A - B\|_1 \leq \frac{1}{2}(C_{op})^{-1} \quad \Rightarrow \quad \|B^{-1}\|_{op} \leq 2C_{op}.$$

Proof. We will use the results stated in Theorem 4.29 of Schott [2016] and the associated discussion in Example 4.14, which establish the following result. Let A be a square, nonsingular matrix, and let I be the identity matrix of the same size. Let $\|\cdot\|$ denote any matrix norm satisfying $\|I\| = 1$. Let D be a matrix of the same size as A satisfying

$$\|A^{-1}\| \|D\| \leq 1. \quad (2)$$

Then

$$\|A^{-1} - (A + D)^{-1}\| \leq \frac{\|A^{-1}\| \|D\|}{1 - \|A^{-1}\| \|D\|} \|A^{-1}\|. \quad (3)$$

We will apply equation (3) using the operator norm $\|\cdot\|_{op}$, for which $\|I\|_{op} = 1$ and with $D := B - A$. Because $\|A^{-1}\|_{op} \leq C_{op}$, A is invertible.

Assume that $\|A - B\|_2 \leq r(C_{op})^{-1}$. First, note that

$$\begin{aligned} \|A^{-1}\|_{op} \|D\|_{op} &= \|A^{-1}\|_{op} \|B - A\|_{op} \\ &\leq \|A^{-1}\|_{op} \|B - A\|_1 \quad (\text{ordering of matrix norms}) \\ &\leq C_{op} \frac{1}{2} (C_{op})^{-1} \quad (\text{by assumption}) \\ &= \frac{1}{2} < 1, \end{aligned} \quad (4)$$

so equation (2) is satisfied and we can apply equation (3). Then

$$\begin{aligned} \|B^{-1}\|_{op} &\leq \|B^{-1} - A^{-1}\|_{op} + \|A^{-1}\|_{op} \quad (\text{triangle inequality}) \\ &\leq \frac{\|A^{-1}\|_{op} \|B - A\|_{op}}{1 - \|A^{-1}\|_{op} \|B - A\|_{op}} \|A^{-1}\|_{op} + \|A^{-1}\|_{op} \quad (\text{equation (3)}) \\ &\leq \frac{\frac{1}{2}}{1 - \frac{1}{2}} \|A^{-1}\|_{op} + \|A^{-1}\|_{op} \quad (\text{equation (4)}) \\ &\leq 2C_{op}. \quad (\text{by assumption}) \end{aligned}$$

□

Next, we define the quantities needed to make use of the integral form of the Taylor series remainder.³

Proposition 3. *For any $\theta \in \Omega_\theta$ and any $w \in W$,*

$$G(\theta, w) - G(\hat{\theta}_1, w) = \left(\int_0^1 H(\hat{\theta}_1 + t(\theta - \hat{\theta}_1), w) dt \right) (\theta - \hat{\theta}_1)$$

Proof. Let $G_d(\theta, w)$ denote the d -th component of the vector $G(\theta, w)$, and define the function $f_d(t) := G_d(\hat{\theta}_1 + t(\theta - \hat{\theta}_1), w)$. The proposition follows by taking the integral remainder form of the zero-th order Taylor series expansion of $f_d(t)$ around $t = 0$ [Dudley, 2018, Appendix B.2], and stacking the result into a vector. \square

The Taylor series residual of Proposition 3 will show up repeatedly, so we will give it a concise name and prove that the integral inherits bounds on the operator norm of the inverse of the Hessian.

Definition 4. For a fixed weight, w , define the Hessian integral

$$\tilde{H}(\theta, w) := \int_0^1 H(\hat{\theta}_1 + t(\theta - \hat{\theta}_1), w) dt.$$

Lemma 2. *If, for some domain D and some constant C , $\sup_{\theta \in D} \|H(\theta, w)^{-1}\|_{op} < C$, then $\sup_{\theta \in D} \|\tilde{H}(\theta, w)^{-1}\|_{op} < C$.*

Proof. By definition of the operator norm,

$$\begin{aligned} \|\tilde{H}(\theta, w)^{-1}\|_{op}^{-1} &= \min_{v \in \mathbb{R}^D: \|v\|_2=1} v^T \tilde{H}(\theta, w) v \\ &= \min_{v \in \mathbb{R}^D: \|v\|_2=1} \int_0^1 v^T H(\hat{\theta}_1 + t(\theta - \hat{\theta}_1), w) v dt \\ &\geq \int_0^1 \min_{v \in \mathbb{R}^D: \|v\|_2=1} v^T H(\hat{\theta}_1 + t(\theta - \hat{\theta}_1), w) v dt \\ &\geq \inf_{\theta \in D} \min_{v \in \mathbb{R}^D: \|v\|_2=1} v^T H(\theta, w) v \\ &\geq C^{-1}. \end{aligned}$$

The result follows. \square

A.2 Lemmas

We now prove some useful consequences of our assumptions. The proof roughly proceeds for all $w \in W_\delta$ by the following steps:

³We are indebted to Pang Wei Koh for pointing out the need to use the integral form of the remainder for Taylor series expansions of vector-valued functions.

1. When δ is small we can make $\left\|\hat{\theta}(w) - \hat{\theta}_1\right\|_2$ small. (Lemma 3 below.)
2. When $\left\|\theta - \hat{\theta}_1\right\|_2$ is small, then the derivatives $H(\theta, w)$ are close to their optimal value, $H(\hat{\theta}_1, 1_w)$. (Lemma 4 and Lemma 5 below.)
3. When the derivatives are close to their optimal values, then $H(\theta, w)$ is uniformly non-singular. (Lemma 6 below.)
4. When the derivatives are close to their optimal values and $H(\theta, w)$ is uniformly non-singular we can control the error in $\hat{\theta}_{\text{IJ}} - \hat{\theta}(w)$ in terms of δ . (Theorem 2 below.)

We begin by showing that the difference between $\hat{\theta}(w)$ and $\hat{\theta}_1$ for $w \in W_\delta$ can be made small by making δ from Condition 1 small.

Lemma 3. *Small parameter changes. Under Assumptions 1—3 and Condition 1,*

$$\text{for all } w \in W_\delta, \quad \left\|\hat{\theta}(w) - \hat{\theta}_1\right\|_2 \leq C_{op}\delta.$$

Proof. By a first-order Taylor expansion in θ , for some $\tilde{\theta}$ such that $\left\|\tilde{\theta} - \hat{\theta}_1\right\|_2 \leq \left\|\hat{\theta}(w) - \hat{\theta}_1\right\|_2$,

$$G(\hat{\theta}(w), 1_w) = G(\hat{\theta}_1, 1_w) + H(\tilde{\theta}, 1_w)(\hat{\theta}(w) - \hat{\theta}_1).$$

By Assumption 2, $H(\tilde{\theta}, 1_w)$ is non-singular. A little manipulation, together with the fact that $G(\hat{\theta}(w), w) = G(\hat{\theta}_1, 1_w) = 0$ gives

$$\begin{aligned} G(\hat{\theta}(w), 1_w) - G(\hat{\theta}(w), w) &= H(\tilde{\theta}, 1_w)(\hat{\theta}(w) - \hat{\theta}_1) \Rightarrow \\ \hat{\theta}(w) - \hat{\theta}_1 &= H(\tilde{\theta}, 1_w)^{-1} \left(G(\hat{\theta}(w), 1_w) - G(\hat{\theta}(w), w) \right). \end{aligned}$$

Applying Condition 1 and Assumption 2,

$$\begin{aligned}
\|\hat{\theta}(w) - \hat{\theta}_1\|_2 &= \left\| H(\tilde{\theta}, 1_w)^{-1} \left(G(\hat{\theta}(w), 1_w) - G(\hat{\theta}(w), w) \right) \right\|_2 \\
&\leq \left\| H(\tilde{\theta}, 1_w)^{-1} \right\|_{op} \left\| G(\hat{\theta}(w), 1_w) - G(\hat{\theta}(w), w) \right\|_2 \\
&\leq \sup_{\theta \in \Omega_\theta} \left\| H(\theta, 1_w)^{-1} \right\|_{op} \left\| G(\hat{\theta}(w), 1_w) - G(\hat{\theta}(w), w) \right\|_2 \\
&\leq C_{op} \left\| G(\hat{\theta}(w), 1_w) - G(\hat{\theta}(w), w) \right\|_2 \quad (\text{Assumption 2}) \\
&\leq C_{op} \left\| G(\hat{\theta}(w), 1_w) - G(\hat{\theta}(w), w) \right\|_1 \quad (\text{relation between norms}) \\
&\leq C_{op} \sup_{\theta \in \Omega_\theta} \|G(\theta, 1_w) - G(\theta, w)\|_1 \\
&\leq C_{op} \delta. \quad (\text{Condition 1}).
\end{aligned}$$

□

Because we will refer to it repeatedly, we give the set of θ defined in Lemma 3 a name.

Definition 5. For a given δ , define the region around $\hat{\theta}_1$ given by Lemma 3 as

$$B_{C_{op}\delta} := \left\{ \theta : \|\theta - \hat{\theta}_1\|_2 \leq C_{op}\delta \right\} \cap \Omega_\theta.$$

In other words, Lemma 3 states that Condition 1 implies $\hat{\theta}(w) \in B_{C_{op}\delta}$ when $w \in W_\delta$.

Next, we show that closeness in θ will mean closeness in $H(\theta, w)$.

Lemma 4. *Boundedness and continuity. Under Assumptions 1–5 and Condition 1,*

$$\text{for all } \theta \in B_{\Delta_\theta}, \quad \sup_{w \in W} \|H(\theta, w) - H(\hat{\theta}_1, w)\|_1 \leq DC_w L_h \|\theta - \hat{\theta}_1\|_2.$$

Proof. For $\theta \in B_{\Delta_\theta}$,

$$\begin{aligned}
\sup_{w \in W} \|H(\theta, w) - H(\hat{\theta}_1, w)\|_1 &= \sup_{w \in W} \left\| \frac{1}{N} \sum_{n=1}^N w_n (h_n(\theta) - h_n(\hat{\theta}_1)) \right\|_1 \quad (\text{by definition}) \\
&\leq D \sup_{w \in W} \frac{\|w\|_2}{\sqrt{N}} \frac{\|h(\theta) - h(\hat{\theta}_1)\|_2}{\sqrt{N}} \quad (\text{Proposition 1}) \\
&\leq DC_w \frac{\|h(\theta) - h(\hat{\theta}_1)\|_2}{\sqrt{N}} \quad (\text{Assumption 5}) \\
&\leq DC_w L_h \|\theta - \hat{\theta}_1\|_2 \quad (\text{Assumption 4 and } \theta \in B_{\Delta_\theta}).
\end{aligned}$$

□

We now combine Lemma 3 and Lemma 4 to show that $H(\theta, w)$ is close to its value at the solution $H(\hat{\theta}_1, 1_w)$ for sufficiently small δ and for all $\theta \in B_{C_{op}\delta}$.

Lemma 5. *Bounds for difference in parameters. Under Assumptions 1–5 and Condition 1, if $\delta \leq \Delta_\theta C_{op}^{-1}$, then*

$$\sup_{\theta \in B_{C_{op}\delta}} \sup_{w \in W_\delta} \left\| H(\theta, w) - H(\hat{\theta}_1, 1_w) \right\|_1 \leq (1 + DC_w L_h C_{op}) \delta.$$

Proof. By Lemma 3, $\delta \leq \Delta_\theta C_{op}^{-1}$ implies that $C_{op}\delta \leq \Delta_\theta$ and so $B_{C_{op}\delta} \subseteq B_{\Delta_\theta}$. Consequently, we can apply Lemma 4:

$$\begin{aligned} \sup_{\theta \in B_{C_{op}\delta}} \sup_{w \in W_\delta} \left\| H(\theta, w) - H(\hat{\theta}_1, w) \right\|_1 &\leq \sup_{\theta \in B_{\Delta_\theta}} \sup_{w \in W_\delta} \left\| H(\theta, w) - H(\hat{\theta}_1, w) \right\|_1 \\ &\leq DC_w L_h \left\| \theta - \hat{\theta}_1 \right\|_2 \quad (\text{Lemma 4}) \\ &\leq DC_w L_h C_{op} \delta \quad (\text{because } \theta \in B_{C_{op}\delta}). \end{aligned}$$

Next, we can use this to write

$$\begin{aligned} \sup_{\theta \in B_{C_{op}\delta}} \sup_{w \in W_\delta} \left\| H(\theta, w) - H(\hat{\theta}_1, 1_w) \right\|_1 &= \sup_{\theta \in B_{C_{op}\delta}} \sup_{w \in W_\delta} \left\| H(\theta, w) - H(\theta, 1_w) + H(\theta, 1_w) - H(\hat{\theta}_1, 1_w) \right\|_1 \\ &\leq \sup_{\theta \in B_{C_{op}\delta}} \sup_{w \in W_\delta} \left\| H(\theta, w) - H(\theta, 1_w) \right\|_1 + \sup_{\theta \in B_{C_{op}\delta}} \sup_{w \in W_\delta} \left\| H(\theta, 1_w) - H(\hat{\theta}_1, 1_w) \right\|_1 \\ &\leq \sup_{\theta \in \Omega_\theta} \sup_{w \in W_\delta} \left\| H(\theta, w) - H(\theta, 1_w) \right\|_1 + \sup_{\theta \in B_{C_{op}\delta}} \sup_{w \in W_\delta} \left\| H(\theta, 1_w) - H(\hat{\theta}_1, 1_w) \right\|_1 \\ &\leq \delta + \sup_{\theta \in B_{C_{op}\delta}} \sup_{w \in W_\delta} \left\| H(\theta, 1_w) - H(\hat{\theta}_1, 1_w) \right\|_1 \quad (\text{Condition 1}) \\ &\leq \delta + DC_w L_h C_{op} \delta. \end{aligned}$$

□

The constant that appears multiplying δ at the end of the proof of Lemma 5 will appear often in what follows, so we give it the special name C_{1J} in Definition 3.

Note that Lemma 5 places a condition on how small δ must be in order for our regularity conditions to apply. Lemma 3 will guarantee that $\hat{\theta}(w) \in B_{C_{op}\delta}$, but if we are not able to make δ arbitrarily small in Condition 1, then we are not guaranteed to ensure that $B_{C_{op}\delta} \subseteq B_{\Delta_\theta}$, will not be able to assume Lipschitz continuity, and none of our results will apply.

Next, using Lemma 5, we can extend the operator bound on H_1^{-1} from Assumption 2 to $H(\theta, w)^{-1}$ for all $w \in W_\delta$, not only for $w = 1_w$.

Lemma 6. *Uniform invertibility of the Hessian. Under Assumptions 1–5 and Condition 1, if $\delta \leq \min \{\Delta_\theta C_{op}^{-1}, \frac{1}{2} C_{IJ}^{-1} C_{op}^{-1}\}$, then*

$$\sup_{\theta \in B_{C_{op}\delta}} \sup_{w \in W_\delta} \|H(\theta, w)^{-1}\|_{op} \leq 2C_{op}.$$

Proof. By Assumption 2, $\|H(\hat{\theta}_1, 1_w)^{-1}\|_{op} \leq C_{op}$. So by Proposition 2, it will suffice to select δ so that

$$\sup_{\theta \in B_{C_{op}\delta}} \sup_{w \in W_\delta} \|H(\theta, w) - H(\hat{\theta}_1, 1_w)\|_1 \leq \frac{1}{2} C_{op}^{-1}. \quad (5)$$

When we can apply Lemma 5, we have

$$\sup_{\theta \in B_{C_{op}\delta}} \sup_{w \in W_\delta} \|H(\theta, w) - H(\hat{\theta}_1, 1_w)\|_1 \leq C_{IJ}\delta.$$

So $H(\theta, w)$ will satisfy equation (5) if we can apply Lemma 5 and if

$$\delta \leq \frac{1}{2} C_{op}^{-1} C_{IJ}^{-1}.$$

To apply Lemma 5 we additionally require that $\delta \leq \Delta_\theta C_{op}^{-1}$. By taking $\delta \leq \min \{\Delta_\theta C_{op}^{-1}, \frac{1}{2} C_{op}^{-1} C_{IJ}^{-1}\}$, we satisfy equation (5) and the result follows. \square

At last, the upper bound on δ will be sufficient to control the error terms in our approximation. For compactness, we give it the name Δ_δ in Definition 3.

Finally, we state a result that will allow us to define derivatives of $\hat{\theta}(w)$ with respect to w .

Lemma 7. *Inverse function theorem. Under Assumptions 1–5 and Condition 1, and for $\delta \leq \Delta_\delta$, there exists a continuous, differentiable function of w , $\hat{\theta}(w)$, such that, for all $w \in W$, $G(\hat{\theta}(w), w) = 0$.*

Proof. This follows from Lemma 6 and the implicit function theorem. \square

By definition, $\hat{\theta}(1_w) = \hat{\theta}_1$.

A.3 Bounding the errors in a Taylor expansion

We are now in a position to use Assumptions 1–5 and Condition 1 to bound the error terms in a first-order Taylor expansion of $\hat{\theta}(w)$. We begin by simply calculating the derivative $d\hat{\theta}(w)/dw$.

Proposition 4. *For any $w \in W$ for which $H(\hat{\theta}(w), w)$ is invertible, and for any vector $a \in \mathbb{R}^N$,*

$$\frac{d\hat{\theta}(w)}{dw^T} \Big|_w a = -H(\hat{\theta}(w), w)^{-1} G(\hat{\theta}(w), a).$$

Proof. Because $G(\hat{\theta}(w), w) = 0$ for all $w \in W$, by direct calculation,

$$\begin{aligned}
0 &= \frac{d}{dw^T} G(\hat{\theta}(w), w)|_w a \\
&= \left(\frac{\partial G}{\partial \theta^T} \frac{d\hat{\theta}}{dw^T} + \frac{\partial G}{\partial w^T} \right) |_w a \\
&= H(\hat{\theta}(w), w) \frac{d\hat{\theta}}{dw^T} |_w a + \left(\frac{\partial}{\partial w^T} \frac{1}{N} \sum_{n=1}^N w_n g_n(\theta) \right) |_w a \\
&= H(\hat{\theta}(w), w) \frac{d\hat{\theta}}{dw^T} |_w a + \frac{1}{N} \sum_{n=1}^N g_n(\hat{\theta}(w)) a \\
&= H(\hat{\theta}(w), w) \frac{d\hat{\theta}}{dw^T} |_w a + G(\hat{\theta}(w), a).
\end{aligned}$$

Because $H(\hat{\theta}(w), w)$ is invertible by assumption, the result follows. \square

Definition 6. Define

$$\begin{aligned}
\hat{\theta}_{IJ}(w) &:= \hat{\theta}_1 + \frac{d\hat{\theta}(w)}{dw^T} |_w (w - 1_w) \\
&= \hat{\theta}_1 - H_1^{-1} G(\hat{\theta}_1, w). \text{ (because } G(\hat{\theta}_1, 1_w) = 0)
\end{aligned}$$

$\hat{\theta}_{IJ}(w)$ in Definition 6 is the first term in a Taylor series expansion of $\hat{\theta}(w)$ as a function of w . We want to bound the error, $\hat{\theta}_{IJ}(w) - \hat{\theta}(w)$.

Theorem 2. Under Assumptions 1-5 and Condition 1, when $\delta \leq \Delta_\delta$,

$$\sup_{w \in W_\delta} \left\| \hat{\theta}_{IJ}(w) - \hat{\theta}(w) \right\|_2 \leq 2C_{op}^2 C_{IJ} \delta^2.$$

Proof. By a one-term Taylor series expansion of $G(\hat{\theta}(w), w) = 0$ in θ around $\hat{\theta}_1$, we have, for some $\tilde{\theta}$ such that $\left\| \tilde{\theta} - \hat{\theta}_1 \right\|_2 \leq \left\| \hat{\theta}(w) - \hat{\theta}_1 \right\|_2$,

$$0 = G(\hat{\theta}(w), w) = G(\hat{\theta}_1, w) + H(\tilde{\theta}, w) (\hat{\theta}(w) - \hat{\theta}_1).$$

Because $\delta \in W_\delta$, Lemma 3 implies that $\hat{\theta}(w) \in B_{C_{op}\delta}$. Because $\left\| \tilde{\theta} - \hat{\theta}_1 \right\|_2 \leq \left\| \hat{\theta}(w) - \hat{\theta}_1 \right\|_2$, $\tilde{\theta} \in B_{C_{op}\delta}$ as well. Because $\tilde{\theta} \in B_{C_{op}\delta}$, Lemma 6 implies that

$H(\tilde{\theta}, w)$ is invertible, so we can solve for $\hat{\theta}(w) - \hat{\theta}_1$.

$$\begin{aligned}\hat{\theta}(w) - \hat{\theta}_1 &= -H(\tilde{\theta}, w)^{-1} G(\hat{\theta}_1, w) \\ &= \left(-H(\tilde{\theta}, w)^{-1} + H(\hat{\theta}_1, 1_w)^{-1} - H(\hat{\theta}_1, 1_w)^{-1} \right) G(\hat{\theta}_1, w) \\ &= \left(H(\hat{\theta}_1, 1_w)^{-1} - H(\tilde{\theta}, w)^{-1} \right) G(\hat{\theta}_1, w) + \hat{\theta}_{\text{IJ}}(w) - \hat{\theta}_1.\end{aligned}$$

Eliminating $\hat{\theta}_1$ and taking the supremum of both sides we have that

$$\begin{aligned}\sup_{w \in W_\delta} \left\| \hat{\theta}_{\text{IJ}}(w) - \hat{\theta}(w) \right\|_2 &= \sup_{w \in W_\delta} \left\| \left(H(\hat{\theta}_1, 1_w)^{-1} - H(\tilde{\theta}, w)^{-1} \right) G(\hat{\theta}_1, w) \right\|_2 \\ &= \sup_{w \in W_\delta} \left\| H(\tilde{\theta}, w)^{-1} \left(H(\tilde{\theta}, w) - H(\hat{\theta}_1, 1_w) \right) H(\hat{\theta}_1, 1_w)^{-1} G(\hat{\theta}_1, w) \right\|_2 \\ &\leq 2C_{op} \sup_{w \in W_\delta} \left\| \left(H(\tilde{\theta}, w) - H(\hat{\theta}_1, 1_w) \right) H(\hat{\theta}_1, 1_w)^{-1} G(\hat{\theta}_1, w) \right\|_2 \quad (\text{Lemma 6}) \\ &\leq 2C_{op} \sup_{w \in W_\delta} \left\| H(\tilde{\theta}, w) - H(\hat{\theta}_1, 1_w) \right\|_{op} \left\| H(\hat{\theta}_1, 1_w)^{-1} G(\hat{\theta}_1, w) \right\|_2 \\ &\leq 2C_{op} \sup_{w \in W_\delta} \left\| H(\tilde{\theta}, w) - H(\hat{\theta}_1, 1_w) \right\|_2 \left\| H(\hat{\theta}_1, 1_w)^{-1} G(\hat{\theta}_1, w) \right\|_2 \quad (\text{ordering of matrix norms}) \\ &\leq 2C_{op} C_{\text{IJ}} \delta \sup_{w \in W_\delta} \left\| H(\hat{\theta}_1, 1_w)^{-1} G(\hat{\theta}_1, w) \right\|_2 \quad (\text{Lemma 5}) \\ &\leq 2C_{op}^2 C_{\text{IJ}} \delta \sup_{w \in W_\delta} \left\| G(\hat{\theta}_1, w) \right\|_2 \quad (\text{Assumption 2}) \\ &= 2C_{op}^2 C_{\text{IJ}} \delta \sup_{w \in W_\delta} \left\| G(\hat{\theta}_1, w) - G(\hat{\theta}_1, 1_w) \right\|_2 \quad (\text{because } G(\hat{\theta}_1, 1_w) = 0) \\ &\leq 2C_{op}^2 C_{\text{IJ}} \delta^2 \quad (\text{Condition 1}).\end{aligned}$$

□

A.4 Use cases

First, let us state a simple condition under which Assumptions 1–4 hold. It will help to have a lemma for the Lipschitz continuity.

Lemma 8. *Derivative Cauchy Schwartz. Let $a(\theta) = (a_1(\theta), \dots, a_N(\theta))$ be an array of tensors with multi-index $i \in [D_A]$, and let $\frac{\partial a(\theta)}{\partial \theta} = (\frac{\partial}{\partial \theta} a_1(\theta), \dots, \frac{\partial}{\partial \theta} a_N(\theta))$ be an array of tensors of size $D \times D_A$. Then*

$$\left\| \frac{\partial}{\partial \theta} \|a(\theta)\|_2 \right\|_2 \leq D_A \left\| \frac{\partial a}{\partial \theta} \right\|_2.$$

Proof. By direct calculation,

$$\begin{aligned}
\left\| \frac{\partial}{\partial \theta} \|a(\theta)\|_2^2 \right\|_2^2 &= \sum_{r=1}^D \left(\frac{\partial}{\partial \theta_r} \sum_{n=1}^N \sum_{i=1}^{D_A} a_{n,i}(\theta)^2 \right)^2 \\
&= \sum_{r=1}^D \left(\sum_{n=1}^N \sum_{i=1}^{D_A} 2a_{n,i}(\theta) \frac{\partial a_{n,i}(\theta)}{\partial \theta_r} \right)^2 \\
&\leq \sum_{r=1}^D \left(2 \sum_{i=1}^{D_A} \left(\sum_{n=1}^N a_{n,i}(\theta)^2 \right)^{\frac{1}{2}} \left(\sum_{n=1}^N \left(\frac{\partial a_{n,i}(\theta)}{\partial \theta_r} \right)^2 \right)^{\frac{1}{2}} \right)^2 \\
&\leq \sum_{r=1}^D \left(2D_A^2 \left(\frac{1}{D_A} \sum_{i=1}^{D_A} \sum_{n=1}^N a_{n,i}(\theta)^2 \right)^{\frac{1}{2}} \left(\frac{1}{D_A} \sum_{n=1}^N \left(\frac{\partial a_{n,i}(\theta)}{\partial \theta_r} \right)^2 \right)^{\frac{1}{2}} \right)^2 \\
&= 4D_A^2 \|a\|_2^2 \sum_{r=1}^D \left\| \frac{\partial a}{\partial \theta_r} \right\|_2^2 \\
&= 4D_A^2 \|a\|_2^2 \left\| \frac{\partial a}{\partial \theta} \right\|_2^2.
\end{aligned}$$

By the chain rule,

$$\left\| \frac{\partial}{\partial \theta} \|a(\theta)\|_2 \right\|_2^2 = \frac{1}{4 \|a(\theta)\|_2^2} \left\| \frac{\partial}{\partial \theta} \|a(\theta)\|_2^2 \right\|_2^2 \leq D_A^2 \left\| \frac{\partial a}{\partial \theta} \right\|_2^2.$$

□

Lemma 9. Let $a(\theta) \in \mathbb{R}^{D \times D}$ be a continuously differentiable random matrix with a $D \times D \times D$ derivative tensor. (Note that the function, not θ , is random. For example, $\mathbb{E}[a(\theta)]$ is still a function of θ .) Suppose that $\mathbb{E}[\|a(\theta)\|_2]$ is finite for all $\theta \in \Omega_\theta$. Then, for all $\theta_1, \theta_2 \in \Omega_\theta$,

$$|\mathbb{E}[\|a(\theta_1)\|_2] - \mathbb{E}[\|a(\theta_2)\|_2]| \leq \sqrt{\mathbb{E} \left[\sup_{\theta \in \Omega_\theta} \left\| \frac{\partial a(\theta)}{\partial \theta} \right\|_2^2 \right]} \|\theta_1 - \theta_2\|_2.$$

Proof. For any tensor a with multi-index i ,

$$\begin{aligned}
\left\| \frac{\partial}{\partial \theta} \|a\|_2^2 \right\|_2^2 &= \sum_{r=1}^D \left(\frac{\partial}{\partial \theta_r} \|a\|_2^2 \right)^2 \\
&= \sum_{r=1}^D \left(\frac{\partial}{\partial \theta_r} \sum_{i=1}^{D_A} a_i^2 \right)^2 \\
&= \sum_{r=1}^D \left(2 \sum_{i=1}^{D_A} a_i \frac{\partial a_i}{\partial \theta_r} \right)^2 \\
&\leq 4 \sum_{r=1}^D \sum_{i=1}^{D_A} a_i^2 \sum_{i=1}^{D_A} \left(\frac{\partial a_i}{\partial \theta_r} \right)^2 \quad (\text{Cauchy-Schwartz}) \\
&= 4 \sum_{i=1}^{D_A} a_i^2 \sum_{r=1}^D \sum_{i=1}^{D_A} \left(\frac{\partial a_i}{\partial \theta_r} \right)^2 \\
&= 4 \|a\|_2^2 \left\| \frac{\partial a}{\partial \theta} \right\|_2^2.
\end{aligned}$$

Consequently,

$$\begin{aligned}
\left\| \frac{\partial}{\partial \theta} \|a(\theta)\|_2 \right\|_2^2 &= \left\| \frac{1}{2 \|a(\theta)\|_2} \frac{\partial}{\partial \theta} \|a(\theta)\|_2^2 \right\|_2^2 \\
&= \frac{1}{4 \|a(\theta)\|_2^2} \left\| \frac{\partial}{\partial \theta} \|a(\theta)\|_2^2 \right\|_2^2 \\
&\leq \frac{4 \|a(\theta)\|_2^2}{4 \|a(\theta)\|_2^2} \left\| \frac{\partial}{\partial \theta} a(\theta) \right\|_2^2 \\
&= \left\| \frac{\partial a(\theta)}{\partial \theta} \right\|_2^2.
\end{aligned}$$

So for any $\theta_1, \theta_2 \in \Omega_\theta$,

$$\begin{aligned}
|\mathbb{E}[\|a(\theta_1)\|_2] - \mathbb{E}[\|a(\theta_2)\|_2]| &\leq \mathbb{E}[|\|a(\theta_1)\|_2 - \|a(\theta_2)\|_2|] \\
&\leq \mathbb{E} \left[\left(\sup_{\theta \in \Omega_\theta} \left\| \frac{\partial}{\partial \theta} \|a(\theta)\|_2 \right\|_2 \right) \right] \|\theta_1 - \theta_2\|_2 \quad (\theta \text{ is not random}) \\
&\leq \mathbb{E} \left[\left(\sup_{\theta \in \Omega_\theta} \left\| \frac{\partial a(\theta)}{\partial \theta} \right\|_2 \right) \right] \|\theta_1 - \theta_2\|_2 \\
&\leq \sqrt{\mathbb{E} \left[\sup_{\theta \in \Omega_\theta} \left\| \frac{\partial a(\theta)}{\partial \theta} \right\|_2^2 \right]} \|\theta_1 - \theta_2\|_2.
\end{aligned}$$

The result follows. Note that the bound still holds (though vacuously) if

$$\mathbb{E} \left[\sup_{\theta \in \Omega_\theta} \left\| \frac{\partial a(\theta)}{\partial \theta} \right\|_2^2 \right] \text{ is infinite.} \quad \square$$

Proposition 5. Let Ω_θ be a compact set. Let $g_n(\theta)$ be twice continuously differentiable IID random functions. Define

$$h_n(\theta) := \frac{\partial g_n(\theta)}{\partial \theta}$$

$$r_n(\theta) := \frac{\partial^2 g_n(\theta)}{\partial \theta \partial \theta},$$

where $r_n(\theta)$ is a $D \times D \times D$ tensor. Assume that

- 1a) $\mathbb{E} \left[\sup_{\theta \in \Omega_\theta} \|g_n(\theta)\|_2^2 \right] < \infty$;
 - 1b) $\mathbb{E} \left[\sup_{\theta \in \Omega_\theta} \|h_n(\theta)\|_2^2 \right] < \infty$;
 - 1c) $\mathbb{E} \left[\sup_{\theta \in \Omega_\theta} \|r_n(\theta)\|_2^2 \right] < \infty$;
 - 2) $\mathbb{E}[h_n(\theta)]$ is non-singular for all $\theta \in \Omega_\theta$;
 - 3) We can exchange expectation and differentiation.
- Then $\lim_{N \rightarrow \infty} P(\text{Assumptions 1-4 hold}) = 1$.

Proof. The proof follows from Theorems 9.1 and 9.2 of Keener [2011]. We will first show that the expected values of the needed functions satisfy Assumptions 1–4, and then that the sample versions converge uniformly.

By Jensen's inequality,

$$\mathbb{E} \left[\sup_{\theta \in \Omega_\theta} \|g_n(\theta)\|_2 \right] = \mathbb{E} \left[\sqrt{\sup_{\theta \in \Omega_\theta} \|g_n(\theta)\|_2^2} \right] \leq \sqrt{\mathbb{E} \left[\sup_{\theta \in \Omega_\theta} \|g_n(\theta)\|_2^2 \right]}.$$

Also, for the i^{th} component of $g_n(\theta)$

$$\mathbb{E} \left[\sup_{\theta \in \Omega_\theta} |g_{n,i}(\theta)| \right] \leq \mathbb{E} \left[\sup_{\theta \in \Omega_\theta} \|g_n(\theta)\|_\infty \right] \leq \mathbb{E} \left[\sup_{\theta \in \Omega_\theta} \|g_n(\theta)\|_2 \right].$$

By Theorem 9.1 of Keener [2011], $\mathbb{E} \left[\|g_n(\theta)\|_2^2 \right]$, $\mathbb{E} [\|g_n(\theta)\|_2]$, and $\mathbb{E} [g_n(\theta)]$ are continuous functions of θ , and because Ω_θ is compact, they are each bounded. Similar reasoning applies to $h_n(\theta)$ and $r_n(\theta)$. Consequently we can define

$$\sup_{\theta \in \Omega_\theta} \mathbb{E} \left[\|g_n(\theta)\|_2^2 \right] =: Q_g^2 < \infty$$

$$\sup_{\theta \in \Omega_\theta} \mathbb{E} \left[\|h_n(\theta)\|_2^2 \right] =: Q_h^2 < \infty.$$

Below, these constants will be used to satisfy Assumption 1 and Assumption 3 with high probability.

Because Ω_θ is compact, $\mathbb{E}[h_n(\theta)]$ is continuous, $\mathbb{E}[h_n(\theta)]$ is non-singular, and the operator norm is a continuous function of $\mathbb{E}[h_n(\theta)]$, we can also define

$$\sup_{\theta \in \Omega_\theta} \left\| \mathbb{E}[h_n(\theta)]^{-1} \right\|_{op} =: Q_{op} < \infty.$$

Below, this constant be used to satisfy Assumption 2 with high probability.

Finally, we turn to the Lipschitz condition. Lemma 9 implies that

$$|\mathbb{E}[\|h_n(\theta_1)\|_2] - \mathbb{E}[\|h_n(\theta_2)\|_2]| \leq \sqrt{\mathbb{E}\left[\sup_{\theta \in \Omega_\theta} \|r_n(\theta)\|_2^2\right]} \|\theta_1 - \theta_2\|_2.$$

Define

$$\Lambda_h = \sqrt{\mathbb{E}\left[\sup_{\theta \in \Omega_\theta} \|r_n(\theta)\|_2^2\right]},$$

so that we have shown that $\mathbb{E}[\|h_n(\theta)\|_2]$ is Lipschitz in Ω_θ with constant Λ_h , which is finite by assumption.

We have now shown, essentially, that the expected versions of the quantities we wish to control satisfy Assumptions 1–4 with $N = 1$. We now need to show that the sample versions satisfy Assumptions 1–4 with high probability, which will follow from the fact that the sample versions converge uniformly to their expectations by Theorem 9.2 of Keener [2011].

First, observe that Assumption 1 holds with probability one by assumption. For the remaining assumption choose an $\epsilon > 0$, and define

$$\begin{aligned} C_g &:= \sqrt{Q_g^2 + \epsilon} \\ C_h &:= \sqrt{Q_h^2 + \epsilon} \\ C_{op} &:= 2Q_{op} \\ L_h &:= \sqrt{D^4 \Lambda_h^2 + \epsilon}. \end{aligned}$$

By Keener [2011] Theorem 9.2,

$$\sup_{\theta \in \Omega_\theta} \left| \frac{1}{N} \sum_{n=1}^N \|g_n(\theta)\|_2^2 - \mathbb{E}[\|g_n(\theta)\|_2^2] \right| \xrightarrow[N \rightarrow \infty]{p} 0.$$

Because

$$\begin{aligned} \sup_{\theta \in \Omega_\theta} \left| \frac{1}{N} \sum_{n=1}^N \|g_n(\theta)\|_2^2 \right| &> Q_g^2 + \epsilon \geq \sup_{\theta \in \Omega_\theta} \mathbb{E}[\|g_n(\theta)\|_2^2] + \epsilon \Rightarrow \\ \sup_{\theta \in \Omega_\theta} \left| \frac{1}{N} \sum_{n=1}^N \|g_n(\theta)\|_2^2 - \mathbb{E}[\|g_n(\theta)\|_2^2] \right| &> \epsilon, \end{aligned}$$

we have

$$\begin{aligned} P\left(\sup_{\theta \in \Omega_\theta} \left| \frac{1}{N} \sum_{n=1}^N \|g_n(\theta)\|_2^2 \right| \geq Q_g^2 + \epsilon\right) &\leq \\ P\left(\sup_{\theta \in \Omega_\theta} \left| \frac{1}{N} \sum_{n=1}^N \|g_n(\theta)\|_2^2 - \mathbb{E}[\|g_n(\theta)\|_2^2] \right| \leq \epsilon\right), \end{aligned}$$

so

$$P \left(\sup_{\theta \in \Omega_\theta} \left| \frac{1}{N} \sum_{n=1}^N \|g_n(\theta)\|_2^2 \right| \geq C_g^2 \right) \xrightarrow{N \rightarrow \infty} 0.$$

An analogous argument holds for $\frac{1}{N} \|h_n(\theta)\|_2^2$. Consequently, $P(\text{Assumption 3 holds}) \xrightarrow{N \rightarrow \infty} 1$.

We now consider Assumption 2. Again, by Keener [2011] Theorem 9.2 applied to each element of the matrix $h_n(\theta)$, using a union bound over each of the D^2 entries,

$$\sup_{\theta \in \Omega_\theta} \left\| \frac{1}{N} \sum_{n=1}^N h_n(\theta) - \mathbb{E}[h_n(\theta)] \right\|_1 \xrightarrow[N \rightarrow \infty]{p} 0.$$

By the converse of Proposition 2, because $\left\| \mathbb{E}[h_n(\theta)]^{-1} \right\|_{op} \leq Q_{op}$,

$$\begin{aligned} \left\| \left(\frac{1}{N} \sum_{n=1}^N h_n(\theta) \right)^{-1} \right\|_{op} &> 2Q_{op} = C_{op} \Rightarrow \\ \left\| \frac{1}{N} \sum_{n=1}^N h_n(\theta) - \mathbb{E}[h_n(\theta)] \right\|_1 &> \frac{1}{2} Q_{op}^{-1}. \end{aligned}$$

Consequently,

$$\begin{aligned} P \left(\left\| \left(\frac{1}{N} \sum_{n=1}^N h_n(\theta) \right)^{-1} \right\|_{op} \geq C_{op} \right) &\leq \\ P \left(\left\| \frac{1}{N} \sum_{n=1}^N h_n(\theta) - \mathbb{E}[h_n(\theta)] \right\|_1 \right) &\xrightarrow[N \rightarrow \infty]{p} 0, \end{aligned}$$

and $P(\text{Assumption 2 holds}) \xrightarrow{N \rightarrow \infty} 1$.

Finally, applying Lemma 9 to $\frac{1}{\sqrt{N}} \|h(\theta_2)\|_2$,

$$\begin{aligned} \left| \frac{1}{\sqrt{N}} \|h(\theta_1)\|_2 - \frac{1}{\sqrt{N}} \|h(\theta_2)\|_2 \right| &\leq \sup_{\theta \in \Omega_\theta} \left\| \frac{\partial}{\partial \theta} \frac{1}{\sqrt{N}} \|h(\theta)\|_2 \right\|_2 \|\theta_1 - \theta_2\|_2 \\ &\leq \frac{D^2}{\sqrt{N}} \sup_{\theta \in \Omega_\theta} \|r(\theta)\|_2 \|\theta_1 - \theta_2\|_2 \\ &= D^2 \sqrt{\sup_{\theta \in \Omega_\theta} \frac{1}{N} \|r(\theta)\|_2^2} \|\theta_1 - \theta_2\|_2. \end{aligned}$$

Consequently,

$$\begin{aligned}
\left| \frac{1}{\sqrt{N}} \|h(\theta_1)\|_2 - \frac{1}{\sqrt{N}} \|h(\theta_2)\|_2 \right| &\geq L_h \|\theta_1 - \theta_2\|_2 \Rightarrow \\
D^2 \sqrt{\sup_{\theta \in \Omega_\theta} \frac{1}{N} \|r(\theta)\|_2^2} &\geq L_h \Rightarrow \\
\sup_{\theta \in \Omega_\theta} \frac{1}{N} \|r(\theta)\|_2^2 - \sup_{\theta \in \Omega_\theta} \mathbb{E} [\|r_n(\theta)\|_2^2] &\geq \frac{L_h^2}{D^4} - \sup_{\theta \in \Omega_\theta} \mathbb{E} [\|r_n(\theta)\|_2^2] \Rightarrow \\
\sup_{\theta \in \Omega_\theta} \left| \frac{1}{N} \|r(\theta)\|_2^2 - \mathbb{E} [\|r_n(\theta)\|_2^2] \right| &\geq \frac{L_h^2}{D^4} - \Lambda_h^2 = \epsilon.
\end{aligned}$$

However, again by Keener [2011] Theorem 9.2,

$$\sup_{\theta \in \Omega_\theta} \left| \frac{1}{N} \|r(\theta)\|_2^2 - \mathbb{E} [\|r_n(\theta)\|_2^2] \right| \xrightarrow[N \rightarrow \infty]{p} 0,$$

so P (Assumption 4 holds) $\xrightarrow[N \rightarrow \infty]{} 1$. □

B Genomics Experiments Details

We demonstrate the Python and R code used to run and analyze the experiments on the genomics data in a sequence of Jupyter notebooks. The output of these notebooks are included below, though they are best viewed in their original notebook form. The notebooks, as well as scripts and instructions for reproducing our analysis in its entirety, can be found in the git repository [rjiordan/AISTATS2019SwissArmyIJ](https://github.com/rjiordan/AISTATS2019SwissArmyIJ).

fit_model_and_save

February 21, 2019

1 Genomics experiment details.

We demonstrate the infinitesimal jackknife on a publicly available data set of mice gene expression in Shoemaker et al. [2015].

Mice were infected with influenza virus, and gene expression was assessed several times after infection, so the observed data consists of expression levels y_{gt} for genes $g = 1, \dots, n_g$ and time points $t = 1, \dots, n_t$, where in this case $n_g = 1000$ and $n_t = 42$.

This notebook contains the first of three steps in the analysis. In this notebook, we will first load the data and define a basis with a hyperparameter we wish to select with cross validation. We then describe the two stages of our analysis: a regression stage and a clustering stage. We then save the data for further analysis by the notebooks `load_and_refit` and `calculate_prediction_error`.

This notebook assumes you have already followed the instructions in `README.md` to install the necessary packages and create the dataset.

2 Step 1: Initial fit.

```
In [1]: import matplotlib.pyplot as plt
        %matplotlib inline

        import numpy as np
        import inspect
        import os
        import sys
        import time

        np.random.seed(3452453) # nothing special about this seed (we hope)!

In [2]: from aistats2019_ij_paper import regression_mixture_lib as rm_lib
        from aistats2019_ij_paper import regression_lib as reg_lib
        from aistats2019_ij_paper import sensitivity_lib as sens_lib
        from aistats2019_ij_paper import spline_bases_lib
        from aistats2019_ij_paper import transform_regression_lib as trans_reg_lib
        from aistats2019_ij_paper import loading_data_utils
        from aistats2019_ij_paper import saving_gmm_utils
        from aistats2019_ij_paper import mse_utils

        import plot_utils_lib
```

2.1 The first stage: Regression

2.1.1 Load data

```
In [3]: # Set bnp_data_repo to be the location of a clone of the repo
# https://github.com/NelleV/genomic_time_series_bnp
bnp_data_repo = '../..//genomic_time_series_bnp'
y_train, y_test, train_indx, timepoints = loading_data_utils.load_genomics_data(
    bnp_data_repo,
    split_test_train = True,
    train_indx_file = '../fits/train_indx.npy')
```

Loading data from: ../..//genomic_time_series_bnp/data/shoemaker2015reprocessed

```
In [4]: n_train = np.shape(y_train)[0]
print('number of genes in training set: \n', n_train)

n_test = np.shape(y_test)[0]
print('number of genes in test set: \n', n_test)

n_genes = n_train + n_test

test_indx = np.setdiff1d(np.arange(n_genes), train_indx)
gene_indx = np.concatenate((train_indx, test_indx))
```

```
number of genes in training set:
700
number of genes in test set:
300
```

Each gene y_g has 42 observations. Observations are made at 14 timepoints, with 3 replicates at each timepoints.

```
In [5]: n_t = len(timepoints)
n_t_unique = len(np.unique(timepoints))

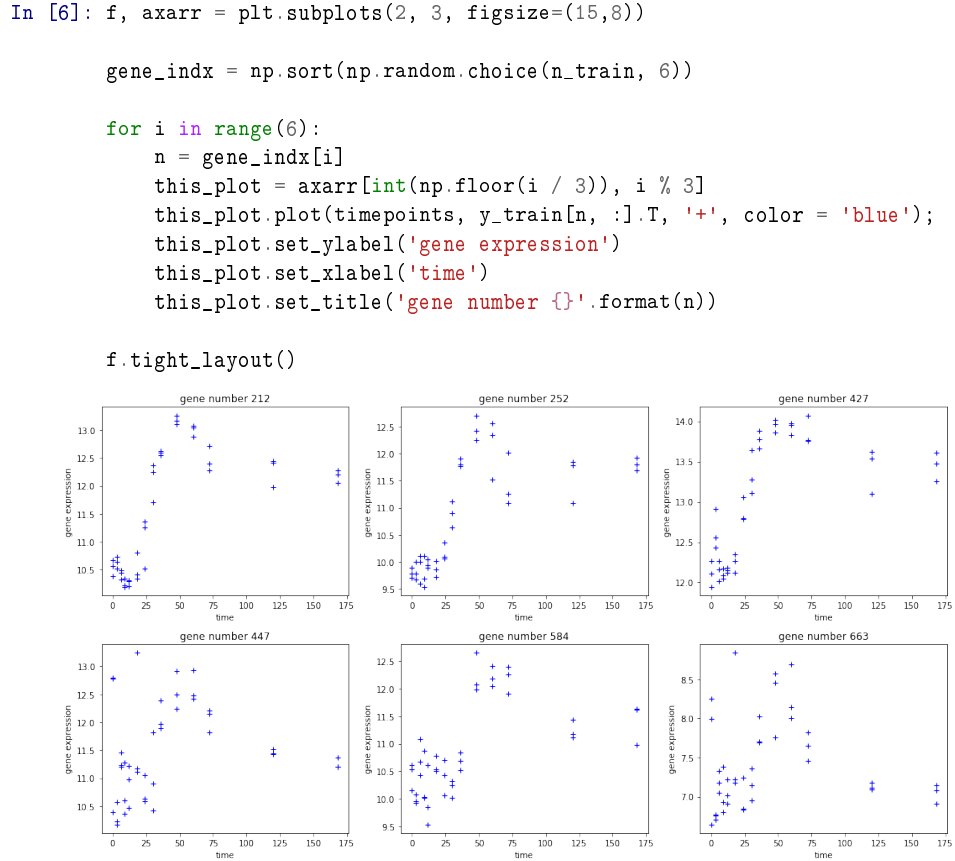
print('timepoints: \n ', timepoints, '\n')
print('Distinct timepoints: \n', np.sort(np.unique(timepoints)), '\n')
print('Number of distinct timepoints:', n_t_unique)
```

```
timepoints:
[0, 0, 0, 3, 3, 3, 6, 6, 6, 9, 9, 9, 12, 12, 12, 18, 18, 18, 24, 24, 24, 30, 30, 30, 36, 36, 3
```

```
Distinct timepoints:
[ 0  3  6  9 12 18 24 30 36 48 60 72 120 168]
```

```
Number of distinct timepoints: 14
```

Here is the raw data for a few randomly chosen genes.



2.1.2 Define regressors

We model the time course using cubic B-splines. Let α be the degrees of freedom of the B-splines, and this is the parameter we seek to choose using cross-validation.

For a given degrees of freedom, the B-spline basis is given by an $n_t \times n_x$ matrix X_{df} , where the each column of X_{df} is a B-spline basis vector evaluated at the n_t timepoints. Note that n_x increases with increasing degrees of freedom.

Note that we only use B-splines to smooth the first 11 timepoints. For the last three timepoints, $t = 72, 120, 168$, we use indicator functions on each timepoint as three extra basis vectors. In other words, we append to the regressor matrix three columns, where each column is 1 if $t = 72, 120$, or 168 , respectively, and 0 otherwise. We do this to avoid numerical issues in the matrix $X^T X$. Because the later timepoints are more spread out, the B-spline basis are close to zero at the later timepoints, leading to matrices close to being singular.

```

In [7]: # Simulate passing arguments in on the command line so that the notebook
        # looks more like those in ``cluster_scripts``.
        class Args():
            def __init__(self):
                pass

        args = Args()
        args.df = 7
        args.degree = 3
        args.num_components = 10

In [8]: regressors = spline_bases_lib.get_genomics_spline_basis(
        timepoints, df=args.df, degree=3)

        regs = reg_lib.Regressions(y_train, regressors)

```

We plot the B-spline matrix for several degrees of freedom below:

```

In [9]: f, axarr = plt.subplots(2, 3, figsize=(15,8))

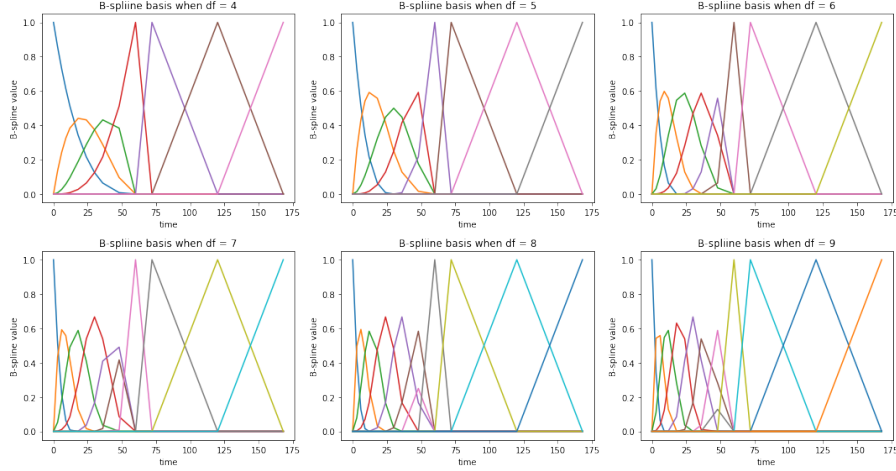
        i = 0
        for df in [4, 5, 6, 7, 8, 9]:
            _regressors = spline_bases_lib.get_genomics_spline_basis(
                timepoints, exclude_num=3, df=df)

            this_plot = axarr[int(np.floor(i / 3)), i % 3]
            this_plot.plot(timepoints, _regressors);
            this_plot.set_xlabel('time')
            this_plot.set_ylabel('B-spline value')
            this_plot.set_title('B-spline basis when df = {}'.format(df))

            i += 1


        f.tight_layout()

```



We display the regressor matrix below.

```
In [10]: plt.matshow(regs.x.T)
plt.ylabel('basis')
plt.xlabel('timepoint and replicate')
plt.title('The (transposed) regressor matrix when df = {}'.format(args.df));
```



With the regressor X defined above, for each gene g we model $P(y_g | \beta_g, \sigma_g^2) = \mathcal{N}(y_g | X\beta_g, \sigma_g^2)$. In the second stage, we will want to cluster β_g taking into account its uncertainty on each gene. To do this, we wish to estimate the posterior mean $\mathbb{E}[\beta_g | y_g]$ and covariance $\text{Cov}(\beta_g | y_g)$ with flat priors for both β_g and σ_g^2 .

For each gene, we estimate the posterior with a mean field variational Bayes (MFVB) approximation $q(\sigma_g^2, \beta_g; \hat{\eta}_g)$ to the posterior $P(\beta_g, \sigma_g^2 | y_g)$.

In particular, we take $q(\sigma_g^2, \beta_g; \hat{\eta}_g) = q^*(\sigma_g^2) q^*(\beta_g)$, where $q^*(\sigma_g^2)$ is a dirac delta function, and we optimize over its a location parameter; $q^*(\beta_g)$ is a Gaussian density and we optimize over its mean and covariance.

The optimal variational approximation has a closed form that is formally identical to the standard frequentist mean and covariance estimate for linear regression. Explicitly, the optimal variational distribution is,

$$q^*(\beta_g) = \mathcal{N}(\beta_g \mid (X^T X)^{-1} X^T y_g, \hat{\tau}_g (X^T X)^{-1})$$

$$q^*(\sigma_g^2) = \delta\{\sigma_g^2 = \hat{\tau}_g\}$$

where $\hat{\tau}_g = \frac{1}{n_t - n_x} \|y_g - X(X^T X)^{-1} X^T y_g\|_2^2$.

The advantage of the MVFB construction is that $\hat{\eta}_g$ for $g = 1, \dots, n_g$ satisfies set of n_g independent M-estimation objectives, allowing us to apply our infinitesimal jackknife results. Specifically, defining $\theta_{reg} := (\eta_1, \dots, \eta_{n_g})$, we wish to minimize

$$F_{reg}(\theta_{reg}, \alpha) = \sum_{g=1}^{n_g} KL(q(\sigma_g^2, \beta_g; \eta_g) \parallel P(\beta_g, \sigma_g^2 | y_g))$$

$$= - \sum_{g=1}^{n_g} \mathbb{E}_q [\log P(\beta_g, \sigma_g^2 | y_g)] + \mathbb{E}_q [\log q(\beta_g, \sigma_g^2 | \eta_g)]$$

$$:= \sum_{g=1}^{n_g} F_{reg,g}(\eta_g, \alpha).$$

Our M-estimator, then, is

$$\frac{\partial F_{reg}(\theta_{reg}, \alpha)}{\partial \theta_{reg}} = 0.$$

The class `regs` can calculate the optimal variational parameters for each gene. In particular, the variational parameters η_g consist of a variational mean and covariance for β_g , as well as a location estimate for σ_g^2 .

```
In [11]: reg_time = time.time()
         opt_reg_params = regs.get_optimal_regression_params()
         reg_time = time.time() - reg_time
         print('Regression time: {} seconds'.format(reg_time))
```

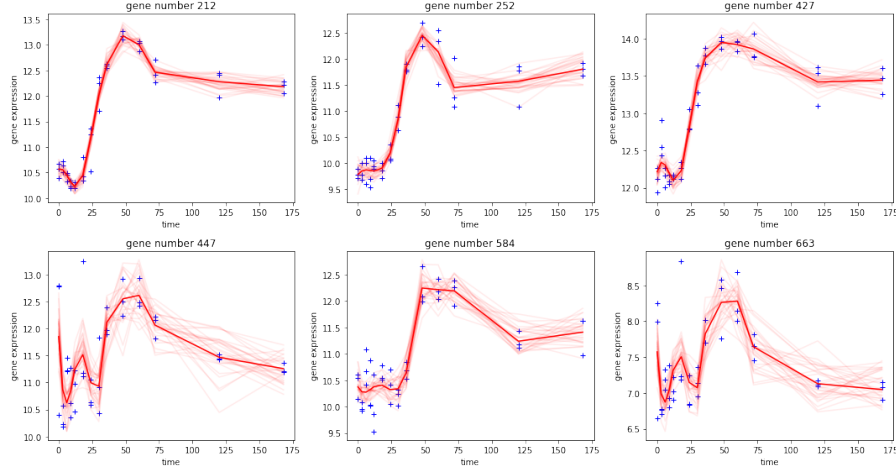
Regression time: 0.029132366180419922 seconds

Here are what some of the fits look like. Each regression produces a prediction $\hat{y}_g := X \mathbb{E}_q[\beta_g]$, plotted with the heavy red line above. The light red are predictions when β_g is drawn from $q^*(\beta_g)$; the spread of the light red is intended to give a sense of the covariance of β_g .

```
In [12]: f, axarr = plt.subplots(2, 3, figsize=(15,8))

         for i in range(6):
             n = gene_indx[i]
             this_plot = axarr[int(np.floor(i / 3)), i % 3]
             plot_utils_lib.PlotRegressionLine(
                 timepoints, regs, opt_reg_params, n, this_plot=this_plot)

         f.tight_layout()
```



We also define and save data for the test regressions, which we will use later to evaluate out-of-sample performance. The training regressions will be saved below with the rest of the fit.

```
In [13]: regs_test = reg_lib.Regressions(y_test, regressors)
         test_regression_outfile = '../fits/test_regressions.json'
         with open(test_regression_outfile, 'w') as outfile:
             outfile.write(regs_test.to_json())
```

2.2 The second stage: fit a mixture model.

2.2.1 Transform the parameters before clustering

We are interested in the pattern of gene expression, not the absolute level, so we wish to cluster $\hat{y}_g - \bar{\hat{y}}_g$, where $\bar{\hat{y}}_g$ is the average over time points. Noting that the $n_t \times n_t$ matrix $\text{Cov}_q(\hat{y}_g - \bar{\hat{y}}_g)$ is rank-deficient because we have subtracted the mean, the final step is to rotate $\hat{y}_g - \bar{\hat{y}}_g$ into a basis where the zero eigenvector is a principle axis and then drop that component.

Call these transformed regression coefficients γ_g and observe that $\text{Cov}_q(\gamma_g)$ has a closed form and is full-rank. It is these γ_g s that we will cluster in the second stage.

We briefly note that the re-centering operation could have been equivalently achieved by making a constant one of the regressors. We chose this implementation because it also allows the user to cluster more complex, non-linear transformations of the regression coefficients, though we leave this extension for future work.

We note that the transformations described in this section are done automatically in the GMM class. We are only calculating these transformations here for exposition.

```
In [14]: # Get the matrix that does the transformation.
         transform_mat, unrotate_transform_mat = \
             trans_reg_lib.get_reversible_predict_and_demean_matrix(regs.x)
         trans_obs_dim = transform_mat.shape[0]
```


If T is the matrix that effects the transformation, then

$$\begin{aligned}\mathbb{E}_q[\gamma_g] &= T\mathbb{E}_q[\beta_g] \\ \text{Cov}_q(\gamma_g) &= T\text{Cov}_q(\beta_g)T^T\end{aligned}$$

The transformed parameters are also regression parameters, just in a different space.

```
In [15]: # Apply the transformation
transformed_reg_params = \
    trans_reg_lib.multiply_regression_by_matrix(
        opt_reg_params, transform_mat)
```

We now visualize the transformed coefficients and their uncertainty.

```
In [16]: f, axarr = plt.subplots(2, 3, figsize=(15,8))


transformed_beta = transformed_reg_params['beta_mean']
transformed_beta_info = transformed_reg_params['beta_info']

for i in range(6):
    n = gene_indx[i]
    this_plot = axarr[int(np.floor(i / 3)), i % 3]
    this_plot.plot(transformed_beta[n, :], color = 'red');
    this_plot.set_ylabel('transformed coefficient')
    this_plot.set_xlabel('index')
    this_plot.set_title('gene number {}'.format(n))

    # draw from the variational distribution, to plot uncertainties
    for j in range(30):
        transformed_beta_draw = np.random.multivariate_normal(
            transformed_beta[n, :], \
            np.linalg.inv(transformed_beta_info[n]))

        axarr[int(np.floor(i / 3)), i % 3].plot(transformed_beta_draw,
            color = 'red', alpha = 0.08);

f.tight_layout()
```



The heavy red lines are the means of the transformed regression coefficients; shaded lines are draws from the variational distribution.

It is these transformed coefficients, γ_g , that we cluster in the second stage.

2.2.2 Estimate an optimal clustering.

We now define a clustering problem for the γ_g . Let n_k be the number of clusters, and μ_1, \dots, μ_{n_k} be the cluster centers. Also let z_{gk} be the binary indicator for the g th gene belonging to cluster k . We then define the following generative model

$$\begin{aligned}
 P(\pi) &= \text{Dirichlet}(\omega) \\
 P(\mu_k) &= \mathcal{N}(\mu_k | 0, \Sigma_0) \quad \text{for } k = 1, \dots, n_k \\
 P(z_{gk} = 1 | \pi_k) &= \pi_k \quad \text{for } k = 1, \dots, n_k; n = 1, \dots, n_g \\
 P(\gamma_g | z_{gk} = 1, \mu_k, \eta_g) &= \mathcal{N}(\gamma_g | \mu_k, \text{Cov}_q(\gamma_g) + \epsilon I_{n_t-1}) \quad \text{for } k = 1, \dots, n_k; n = 1, \dots, n_g.
 \end{aligned}$$

where ϵ is a small regularization parameter, which helped our optimization produce more stable results.

We will estimate the clustering using the maximum a posteriori (MAP) estimator of $\theta_{clust} := (\mu, \pi)$. This defines an optimization objective that we seek to minimize:

$$F_{clust}(\theta_{clust}, \theta_{reg}) = - \sum_{g=1}^{n_g} E_{q_z^*} \left\{ \log P(\gamma_g | \eta_g, \mu, \pi, z_g) - \log P(z_g | \pi) \right\} - \log P(\mu) - \log P(\pi)$$

which, for every value of θ_{reg} , we expect to satisfy

$$\frac{\partial F_{clust}(\theta_{clust}, \theta_{reg})}{\partial \theta_{clust}} = 0.$$

Note that θ_{clust} involves only the “global” parameters μ and π . We did take a variational distribution for the z_{gk} s, represented by independent Bernoulli distribution, but the optimal q_z^* can be written as a function of μ and π . Hence, our optimization objective only involves these global parameters.

```
In [17]: # Define prior parameters.
         num_components = args.num_components
         epsilon = 0.1
         loc_prior_info_scalar = 1e-5

         trans_obs_dim = regs.x.shape[1] - 1
         prior_params = \
             rm_lib.get_base_prior_params(trans_obs_dim, num_components)
         prior_params['probs_alpha'][:] = 1
         prior_params['centroid_prior_info'] = loc_prior_info_scalar * np.eye(trans_obs_dim)

In [18]: gmm = rm_lib.GMM(args.num_components,
                           prior_params, regs, opt_reg_params,
                           inflate_coef_cov=None,
                           cov_regularization=epsilon)
```

In our experiment, the number of clusters n_k was chosen to be 10. We set ω to be the ones vector of length n_k . The prior info for the cluster centers Σ_0 is $1e-05 \times I$. ϵ was set to be 0.1.

Let us examine the optimization objective. First, we’ll inspect the likelihood terms. What follows is the likelihood given that gene g belongs to cluster k .

```
In [19]: print(inspect.getsource(rm_lib.get_log_lik_nk))

def get_log_lik_nk(centroids, probs, x, x_infos):
    loc_log_lik = \
        -0.5 * (-2 * np.einsum('ni,kj,nij->nk', x, centroids, x_infos) +
               np.einsum('ki,kj,nij->nk', centroids, centroids, x_infos))

    log_probs = np.log(probs[0, :])
    log_lik_by_nk = loc_log_lik + log_probs.T

    return log_lik_by_nk
```

We can then optimize for q_z^* , which can be parametrized by its mean $\mathbb{E}_{q_z^*}[z]$. We note that this update has a closed form given θ_{clust} , so there is no need to solve an optimization problem to find $q_z^*(z)$. We additionally note that we do not use the EM algorithm, which we found to have exhibit extremely poor convergence rates. Rather, we set $q_z^*(z)$ to its optimal value given θ_{clust} and return the objective as a function of θ_{clust} alone, allowing the use of more general and higher-quality optimization routines.

```
In [20]: print(inspect.getsource(rm_lib.get_e_z))
```

```
def get_e_z(log_lik_by_nk):
    log_const = paragami.simplex_patterns.logsumexp(log_lik_by_nk, axis=1)
    e_z = np.exp(log_lik_by_nk - log_const)
    return e_z
```

With the optimal parameters for z_{nk} , we combine the likelihood term with the prior and entropy terms.

```
In [21]: print(inspect.getsource(rm_lib.wrap_get_loglik_terms))
         print(inspect.getsource(rm_lib.wrap_get_kl))

def wrap_get_loglik_terms(gmm_params, transformed_reg_params):
    log_lik_by_nk = get_log_lik_nk(
        centroids=gmm_params['centroids'],
        probs=gmm_params['probs'],
        x=transformed_reg_params['beta_mean'],
        x_infos=transformed_reg_params['beta_info'])

    e_z = get_e_z(log_lik_by_nk)

    return log_lik_by_nk, e_z

def wrap_get_kl(gmm_params, transformed_reg_params, prior_params):
    log_lik_by_nk, e_z = \
        wrap_get_loglik_terms(gmm_params, transformed_reg_params)
    log_prior = get_log_prior(
        gmm_params['centroids'], gmm_params['probs'], prior_params)
    return get_kl(log_lik_by_nk, e_z, log_prior)
```

This objective function is wrapped in the GMM class method `get_params_kl`.

```
In [22]: print(inspect.getsource(gmm.get_params_kl))

def get_params_kl(self, gmm_params):
    """Get the optimization objective as a function of the mixture
    parameters.
    """
    return wrap_get_kl(
        gmm_params, self.transformed_reg_params, self.prior_params)
```

2.2.3 Optimization

For optimization we make extensive use of the [autograd](#) library for automatic differentiation and the [paragami](#) library for parameter packing and sparse Hessians. These packages' details are beyond the scope of the current notebook.

First, we do a k-means initialization.

```
In [23]: print('Running k-means init.')
         init_gmm_params = \
             rm_lib.kmeans_init(gmm.transformed_reg_params,
                               gmm.num_components, 50)

         print('Done.')
         init_x = gmm.gmm_params_pattern.flatten(init_gmm_params, free=True)
```

```
Running k-means init.
Done.
```

We note that the match between “exact” cross-validation (removing time points and re-optimizing) and the IJ was considerably improved by using a high-quality second-order optimization method. In particular, for these experiments, we employed the Newton conjugate-gradient trust region method (Chapter 7.1 of Wright et al [1999]) as implemented by the method `trust-ncg` in `scipy.optimize`, preconditioned by the Cholesky decomposition of an inverse Hessian calculated at an initial approximate optimum.

We found that first-order or quasi-Newton methods (such as BFGS) often got stuck or terminated at points with fairly large gradients. At such points our method does not apply in theory nor, we found, very well in practice.

The inverse Hessian used for the preconditioner was with respect to the clustering parameters only and so could be calculated quickly, in contrast to the H_1 matrix used for the IJ, which includes the regression parameters as well.

First, run with a low tolerance to get a point at which to evaluate an initial preconditioner.

```
In [24]: gmm.conditioned_obj.reset() # Reset the logging and iteration count.
         gmm.conditioned_obj.set_print_every(1)

         opt_time = time.time()
         gmm_opt, init_x2 = gmm.optimize(init_x, gtol=1e-2)
         opt_time = time.time() - opt_time
```

```
Iter 0: f = -159.11834165
Iter 1: f = -159.67926278
Iter 2: f = -159.97782885
Iter 3: f = -160.15878320
Iter 4: f = -159.59447036
Iter 5: f = -160.19209687
Iter 6: f = -160.27259154
Iter 7: f = -160.29486553
Iter 8: f = -160.33460656
Iter 9: f = -160.34154288
Iter 10: f = -160.32382096
Iter 11: f = -160.34447865
Iter 12: f = -160.34634639
Iter 13: f = -160.34692896
```

Next, set the preconditioner using the square root inverse Hessian at the point `init_x2`.

```
In [25]: tic = time.time()
         h_cond = gmm.update_preconditioner(init_x2)
         opt_time += time.time() - tic
```

The method `optimize_fully` repeats this process of optimizing and re-calculating the preconditioner until the optimal point does not change.

```
In [26]: gmm.conditioned_obj.reset()
         tic = time.time()
         gmm_opt, gmm_opt_x = gmm.optimize_fully(
             init_x2, verbose=True)
         opt_time += time.time() - tic
         print('Optimization time: {} seconds'.format(opt_time))
```

```
Preconditioned iteration 1
  Running preconditioned optimization.
Iter 0: f = -160.34692896
Iter 1: f = -160.34694250
Iter 2: f = -160.34694250
Preconditioned iteration 2
  Getting Hessian and preconditioner.
  Running preconditioned optimization.
Iter 3: f = -160.34694250
Iter 4: f = -160.34694250
Converged.
Optimization time: 8.438910484313965 seconds
```

`paragami` patterns allow conversion between unconstrained vectors and dictionaries of parameter values. After “folding” the optimal `gmm_opt_x`, `opt_gmm_params` contains a dictionary of optimal cluster centroids and cluster probabilities.

```
In [27]: opt_gmm_params = gmm.gmm_params_pattern.fold(gmm_opt_x, free=True)
         print(opt_gmm_params.keys())
         print(np.sort(opt_gmm_params['probs']))

odict_keys(['centroids', 'probs'])
[[0.01567608 0.04016882 0.06955236 0.07427946 0.09373695 0.0947442
  0.09653288 0.12626624 0.15739176 0.23165127]]
```

Each gene’s regression line has an inferred cluster membership given by $\mathbb{E}_{q_z^*}[z_g]$, and an expected posterior centroid given by $\sum_k \mathbb{E}_{q_z^*}[z_{gk}] \mu_k$. This expected posterior centroid can be untransformed to give a prediction for the observation.

It is the difference between this prediction line — which is a function of the clustering — and the actual data that we consider to be the “error” of the model.

```

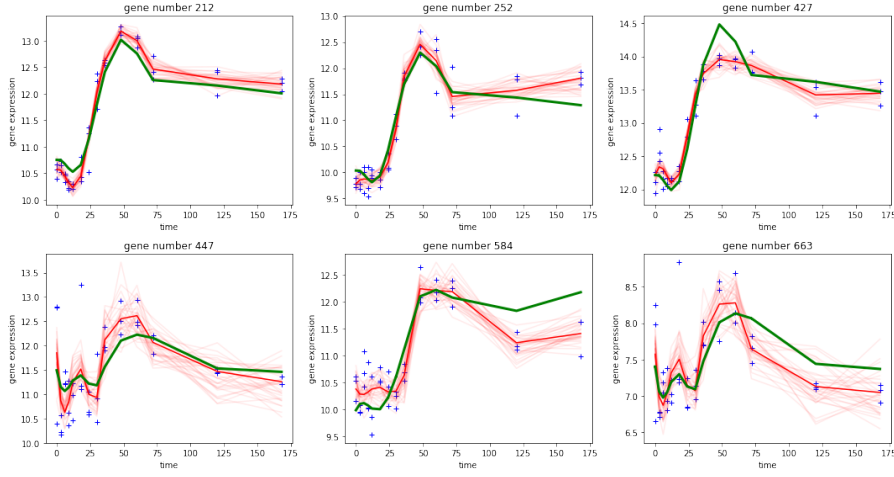
In [28]: gmm_pred = mse_utils.get_predictions(gmm, opt_gmm_params, opt_reg_params)

f, axarr = plt.subplots(2, 3, figsize=(15,8))

for i in range(6):
    n = gene_indx[i]
    this_plot = axarr[int(np.floor(i / 3)), i % 3]
    plot_utils_lib.PlotRegressionLine(
        timepoints, regs, opt_reg_params, n, this_plot=this_plot)
    plot_utils_lib.PlotPredictionLine(
        timepoints, regs, gmm_pred, n, this_plot=this_plot)

f.tight_layout()

```



2.2.4 Calculating H_1 for the IJ

We seek to choose the degrees of freedom α for the B-splines using cross-validation. We leave out one or more timepoints, and fit using only the remaining timepoints. We then estimate the test error by predicting the value of the genes at the held out timepoints.

To do this, we define time weights w_t by observing that, for each g , the term $\mathbb{E}_q \left[\log P \left(\beta_g, \sigma_g^2 | y_g \right) \right]$ decomposes into a sum over time points:

$$F_{reg,g}(\eta_g, \alpha, w) := - \sum_{t=1}^{n_t} w_t \left(-\frac{1}{2} \sigma_g^{-2} (y_{g,t} - (X\beta_g)_t)^2 - \frac{1}{2} \log \sigma_g^2 \right) + \mathbb{E}_q \left[\log q \left(\beta_g, \sigma_g^2 | \eta_g \right) \right].$$

We naturally define $F_{reg}(\theta_{reg}, \alpha, w) := \sum_{g=1}^{n_g} F_{reg,g}(\eta_g, \alpha, w)$.

By defining $\theta = (\theta_{clust}, \theta_{reg})$, we then have an M-estimator

$$G(\theta, w, \alpha) := \begin{pmatrix} \frac{\partial F_{reg}(\theta_{reg}, w, \alpha)}{\partial \theta_{reg}} \\ \frac{\partial F_{clust}(\theta_{clust}, \theta_{reg})}{\partial \theta_{clust}} \end{pmatrix} = 0.$$

We can then apply the IJ to approximate the leaving out of various timepoints.

Note that what we call the “Hessian” for this two-step procedure is not really a Hessian, as it is not symmetric. More precisely, it is the Jacobian of G , or what we defined as H_1 in the text.

Calculating H_1 is the most time-consuming part of the infinitesimal jackknife, since the H_1 matrix is quite large (though sparse). However, once H_1 is computed, calculating each $\theta_{IJ}(w)$ is extremely fast.

H_1 can be computed in blocks:

$$H_1 = \begin{pmatrix} \nabla_{\theta_{reg}}^2 F_{reg} & 0 \\ \nabla_{\theta_{reg}} \nabla_{\theta_{clust}} F_{clust} & \nabla_{\theta_{clust}}^2 F_{clust} \end{pmatrix}$$

The code refers to $\nabla_{\theta_{clust}}^2 F_{clust}$ as the “GMM Hessian”. It refers to $\nabla_{\theta_{reg}} \nabla_{\theta_{clust}} F_{clust}$ as the “cross Hessian”. And it refers to $\nabla_{\theta_{reg}}^2 F_{reg}$ as the “regression Hessian”, which itself is block diagonal, with each block an observation. Due to details of the implementation of block sparse Hessians using forward mode automatic differentiation in the class `vittles.SparseBlockHessian`, the code below confusingly refers to each regression parameter as a “block”.

When the `FitDerivatives` class is initialized, it calculates these blocks separately and stacks them into the attribute `full_hess`, which is a sparse matrix representing H_1 .

```
In [29]: # Even though $H_1$ is not a Hessian, by force of habit we call the time to
# compute it ``hess_time``.
hess_time = time.time()
fit_derivs = sens_lib.FitDerivatives(
    opt_gmm_params, opt_reg_params,
    gmm.gmm_params_pattern, regs.reg_params_pattern,
    gmm=gmm, regs=regs,
    print_every=10)
hess_time = time.time() - hess_time
print('Total hessian time: {} seconds'.format(hess_time))
```

```
Initializing FitDerivatives.
Getting t Jacobian.
Getting full Hessian.
    Getting GMM Hessian...
    GMM Hessian time: 2.1917014122009277
    Getting cross Hessian...
    Cross Hessian time: 34.25235605239868
    Getting regression Hessian...
Block index 0 of 66.
Block index 10 of 66.
Block index 20 of 66.
Block index 30 of 66.
```



```

Block index 40 of 66.
Block index 50 of 66.
Block index 60 of 66.
Done differentiating.
Regression Hessian time: 121.74362897872925
Done with full Hessian.
Total hessian time: 169.59288716316223 seconds

```

2.2.5 Save results as a compressed file.

The results, including H_1 , are now saved. To calculate the exact CV, these results (including the preconditioner) will be loaded and the model will be refit with timepoints left out. To calculate the IJ, the same results will be loaded and H_1 will be used to calculate the IJ.

```

In [30]: extra_metadata = dict()
         extra_metadata['opt_time'] = opt_time
         extra_metadata['reg_time'] = reg_time
         extra_metadata['hess_time'] = hess_time
         extra_metadata['df'] = args.df
         extra_metadata['degree'] = args.degree

         npz_outfile = '../fits/initial_fit.npz'
         saving_gmm_utils.save_initial_optimum(
             npz_outfile,
             gmm=gmm,
             regs=regs,
             timepoints=timepoints,
             fit_derivs=fit_derivs,
             extra_metadata=extra_metadata)

```

2.2.6 Bibliography

J. E. Shoemaker, S. Fukuyama, A. J. Einfeld, D. Zhao, E. Kawakami, S. Sakabe, T. Maemura, T. Gorai, H. Katsura, Y. Muramoto, S. Watanabe, T. Watanabe, K. Fuji, Y. Matsuoka, H. Kitano, and Y. Kawaoka. An Ultrasensitive Mechanism Regulates Influenza Virus-Induced Inflammation. *PLoS Pathogens*, 11(6):1–25, 2015

S. Wright and J. Nocedal. Numerical optimization. Springer Science, 35(67–68):7, 1999.

load_and_refit

February 21, 2019

1 Step 2: Refit.

In this notebook, we calculate the parameters used for exact CV by refitting the model initially fit in step one, the notebook `fit_model_and_save`.

For expository purposes this notebook calculates the refit for only one weight vector. To compute exact CV, one would perform the corresponding computation for all leave-k-out weight vectors.

```
In [1]: from copy import deepcopy
import inspect
import matplotlib.pyplot as plt
%matplotlib inline
import numpy as np
import sys
import time

np.random.seed(3452453)

import paragami

from aistats2019_ij_paper import regression_mixture_lib as rm_lib
from aistats2019_ij_paper import saving_gmm_utils
from aistats2019_ij_paper import mse_utils

import plot_utils_lib
```

```
In [2]: # Load the initial fit.
# This file was produced by the notebook ``fit_model_and_save``.
initial_fit_infile = '../fits/initial_fit.npz'
full_fit, gmm, regs, metadata = \
    saving_gmm_utils.load_initial_optimum(initial_fit_infile)
timepoints = metadata['timepoints']
```

```
Initializing FitDerivatives.
Using provided t_jac.
Using provided full_hess.
```

First, choose some timepoints to leave out.

```
In [3]: # Simulate passing arguments in on the command line.
class Args():
    def __init__(self):
        pass

args = Args()
args.num_times = 1
args.which_comb = 1
args.max_num_timepoints = 7
```

The number of points left out (that is, k) is given by `num_times`, which is 1. The largest time-point we leave out is given by `max_num_timepoints`, which is 7. Because later timepoints are not affected by the smoothing, there is no reason to leave them out.

There are a certain number of ways to leave k out of 7 timepoints, and `which_comb` chooses one of them in the order given by the function `itertools.combinations`. Of course, when $k = 1$, `which_comb` simply chooses which timepoint to leave out. `mse_utils.get_indexed_combination` maps `which_comb` to particular timepoints in a consistent way.

Full exact CV would run this script for all 7 choose k values of which_comb.

Because we have repeated measurements at each timepoint, leaving out a single timepoint will correspond to leaving out multiple row of the observation matrix. Those rows are determined by `mse_utils.get_time_weight`, which also returns a weight vector setting these observations' weights to zero.

```
In [4]: lo_inds = mse_utils.get_indexed_combination(
        num_times=args.num_times, which_comb=args.which_comb,
        max_num_timepoints=args.max_num_timepoints)
        new_time_w, full_lo_inds = mse_utils.get_time_weight(lo_inds, timepoints)

print('Left out timepoint: {}'.format(lo_inds))
print('Left out observations: {}'.format(full_lo_inds))
print('Leave-k-out weights: {}'.format(new_time_w))
```

[illegible]

We now re-optimize with the new weights.

Note that we could either start the optimization at the initial optimum (a “warm start”) or do a fresh start from k-means. A fresh start is more time consuming but a more stringent test for the accuracy of the IJ. We calculate both, but report results from the fresh start in the paper. In the notebook `examine_and_save_results`, you can choose to examine either set of results.

Here, for consistency with the paper, we re-initialize with k-means.

```
In [5]: regs.time_w = deepcopy(new_time_w)
        reg_params_w = regs.get_optimal_regression_params()
```

```

gmm.set_regression_params(reg_params_w)

init_gmm_params = \
    rm_lib.kmeans_init(gmm.transformed_reg_params,
                       gmm.num_components, 50)
init_x = gmm.gmm_params_pattern.flatten(init_gmm_params, free=True)

opt_time = time.time()
gmm_opt, init_x2 = gmm.optimize(init_x, gtol=1e-2)

print('\tUpdating preconditioner...')
kl_hess = gmm.update_preconditioner(init_x2)

print('\tRunning preconditioned optimization...')
gmm.conditioned_obj.reset()
reopt, gmm_params_free_w = gmm.optimize_fully(init_x2, verbose=True)
print(gmm_opt.message)
opt_time = time.time() - opt_time

print('Refit time: {} seconds'.format(opt_time))

Iter 0: f = -153.38003431
Iter 1: f = -152.49438715
Iter 2: f = -153.69147895
Iter 3: f = -153.83779915
Iter 4: f = -154.02397812
Iter 5: f = -153.41393391
Iter 6: f = -154.10396420
Iter 7: f = -154.14366282
Iter 8: f = -154.14261201
Iter 9: f = -154.16417745
Iter 10: f = -154.18307547
Iter 11: f = -154.20711481
Iter 12: f = -154.22118064
Iter 13: f = -154.27402715
Iter 14: f = -154.28739474
Iter 15: f = -154.33849929
Iter 16: f = -154.03580241
Iter 17: f = -154.35421130
Iter 18: f = -154.36910489
Iter 19: f = -154.36872458
Iter 20: f = -154.37238982
Iter 21: f = -154.37722095
Iter 22: f = -154.38186985
Iter 23: f = -154.38410992
    Updating preconditioner...
    Running preconditioned optimization...
Preconditioned iteration 1

```

```

Running preconditioned optimization.
Iter 0: f = -154.38410992
Iter 1: f = -154.38423176
Iter 2: f = -154.38584092
Iter 3: f = -154.21889674
Iter 4: f = -154.42200228
Iter 5: f = -154.39603234
Iter 6: f = -154.39957947
Iter 7: f = -154.41374585
Iter 8: f = -154.43397491
Iter 9: f = -154.43484046
Iter 10: f = -154.43484816
Iter 11: f = -154.43484816
Preconditioned iteration 2
Getting Hessian and preconditioner.
Running preconditioned optimization.
Iter 12: f = -154.43484816
Iter 13: f = -154.43484816
Converged.
Optimization terminated successfully.
Refit time: 14.35115647315979 seconds

```

We now save the results.

```

In [6]: gmm_params_w = \
        full_fit.comb_params_pattern['mix'].fold(
            gmm_params_free_w, free=True)
refit_comb_params = {
    'mix': gmm_params_w,
    'reg': reg_params_w }
refit_comb_params_free = \
    full_fit.comb_params_pattern.flatten(refit_comb_params, free=True)

In [7]: save_filename = \
        './fits/refit__num_times{}___which_comb{}.npz'.format(
            args.num_times, args.which_comb)
print('Saving to {}'.format(save_filename))
saving_gmm_utils.save_refit(
    outfile=save_filename,
    comb_params_free=refit_comb_params_free,
    comb_params_pattern=full_fit.comb_params_pattern,
    initial_fit_infile=initial_fit_infile,
    time_w=new_time_w,
    lo_inds=lo_inds,
    full_lo_inds=full_lo_inds)

Saving to ./fits/refit__num_times1__which_comb1.npz

```

calculate_prediction_errors

February 21, 2019

1 Step 3: Calculate the IJ and prediction errors.

In this notebook, for a single weight vector, we calculate the IJ itself as well as the prediction errors for exact CV and IJ. This notebook uses the output of the notebooks `load_and_refit` and `fit_model_and_save`.

```
In [1]: import numpy as np
import paragami
import vittles
import scipy as sp
from scipy import sparse
import time

import seaborn as sns
import pandas as pd

import matplotlib.pyplot as plt
%matplotlib inline

np.random.seed(3452453)

from aistats2019_ij_paper import regression_lib as reg_lib
from aistats2019_ij_paper import sensitivity_lib as sens_lib
from aistats2019_ij_paper import saving_gmm_utils
from aistats2019_ij_paper import mse_utils

import plot_utils_lib

In [2]: # Simulate passing arguments in on the command line.
class Args():
    def __init__(self):
        pass

args = Args()
args.num_times = 1
args.which_comb = 1
args.max_num_timepoints = 7
```

```

In [3]: #####
        # Load the original fit.

        print('Loading original fit.')
        initial_fit_infile = '../fits/initial_fit.npz'
        full_fit, gmm, regs, initial_metadata = \
            saving_gmm_utils.load_initial_optimum(initial_fit_infile)

        opt_comb_params = full_fit.get_comb_params()

Loading original fit.
Initializing FitDerivatives.
Using provided t_jac.
Using provided full_hess.

In [4]: #####
        # Load the test data

        test_regression_infile = '../fits/test_regressions.json'
        with open(test_regression_infile) as infile:
            regs_test = reg_lib.Regressions.from_json(infile.read())

#####
        # Load a refit as specified by ``args``.

        refit_filename = \
            '../fits/refit__num_times{__which_comb__}.npz'.format(
                args.num_times, args.which_comb)

        comb_params_free_refit, comb_params_pattern_refit, refit_metadata = \
            saving_gmm_utils.load_refit(refit_filename)

        time_w = refit_metadata['time_w']
        lo_inds = refit_metadata['lo_inds']
        full_lo_inds = refit_metadata['full_lo_inds']

        assert(comb_params_pattern_refit == full_fit.comb_params_pattern)
        comb_params_refit = comb_params_pattern_refit.fold(
            comb_params_free_refit, free=True)

        time_w = refit_metadata['time_w']
        lo_inds = refit_metadata['lo_inds']
        full_lo_inds = refit_metadata['full_lo_inds']

```

The objects named `comb_params` refer to both the regression and clustering parameters. The name `free` refers to the unconstrained flat value for the parameters as calculated by `paragami`.

```

In [5]: print('Regression pattern: ',
            comb_params_pattern_refit['reg'])

```

```

print('Clustering pattern: ',
      comb_params_pattern_refit['mix'])

Regression pattern: OrderedDict:
  [beta_mean] = NumericArrayPattern (700, 10) (lb=-inf, ub=inf)
  [beta_info] = PatternArray (700,) of PDMatrix 10x10 (diag_lb = 0.0)
  [y_info] = NumericArrayPattern (700,) (lb=0.0, ub=inf)
Clustering pattern: OrderedDict:
  [centroids] = NumericArrayPattern (10, 9) (lb=-inf, ub=inf)
  [probs] = SimplexArrayPattern (1,) of 10-d simplices

```

1.0.1 Calculate the infinitesimal jackknife.

The vittles package makes it easy to calculate linear approximations to the sensitivity of M-estimators to hyperparameters, of which the IJ is a special case. Here, the HyperparameterSensitivityLinearApproximation uses the sparse value of H_1 calculated earlier.

Note that H_1 is factorized during the initialization of `weight_sens`, and that it takes relatively little time.

```

In [6]: # Note that if you don't cast the jacobian to a numpy array from
# a numpy matrix, the output is a 2d-array, causing confusion later.
weight_sens = vittles.HyperparameterSensitivityLinearApproximation(
    objective_fun=lambda: 0,
    opt_par_value=full_fit.comb_params_free,
    hyper_par_value=regs.time_w,
    hessian_at_opt=sp.sparse.csc_matrix(full_fit.full_hess),
    cross_hess_at_opt=np.array(full_fit.t_jac.todense()))

```

We now use the `weight_sens` object to approximate the “free” value of the combined parameters at `time_w`. The IJ operates in unconstrained space, so we use `paragami` to fold the unconstrained vector back into a dictionary of parameters.

```

In [7]: # Get the infinitesimal jackknife for the refit weight vector.
lr_time = time.time()
comb_params_free_lin = \
    weight_sens.predict_opt_par_from_hyper_par(time_w)
lr_time = time.time() - lr_time
print('Infinitesimal jackknife time: {}'.format(lr_time))

comb_params_lin = full_fit.comb_params_pattern.fold(comb_params_free_lin, free=True)

Infinitesimal jackknife time: 0.0011603832244873047

```

1.0.2 Calculate various prediction errors.

Recall that the prediction error is the difference between the data and the posterior expected cluster centroid for a particular gene. Let us consider the original optimal clustering parameters,

`opt_comb_params['mix']`. To get the test set error on gene g for these parameters, we need to do the following steps:

1. Run the regression for gene g in the test set
2. Classify the regression, calculating $\mathbb{E}_{q_z^*}[z_g]$. This is a function of the clustering parameters and the regression line for gene g .
3. Calculate the expected posterior cluster centroid for gene g , which is $\mu_g^* = \sum_k \mathbb{E}_{q_z^*}[z_{gk}] \mu_k$.
4. Because the transformation discards the mean information, compare the de-means data to the estimated centroid: $error_{gt} = \left(y_{gt} - \frac{1}{T} \sum_{t'=1}^T y_{gt'}\right) - \mu_{gt}^*$.

Note that step one could re-run the regression either with the original weights or the new weights. We found that this decision does not matter qualitatively. Here and in the paper, we simply classify the original regression, but the notebook `examine_and_save_results` can produce results for both the original and re-weighted regressions.

We will examine prediction error on the time points that are left out, that is, for observations in `full_lo_inds`.

```
In [8]: print('Calculating prediction error.')

# Get the training set error on the full data.
train_error = mse_utils.get_lo_err_folded(
    opt_comb_params,
    keep_inds=full_lo_inds,
    mse_regs=regs,
    mse_reg_params=opt_comb_params['reg'],
    gmm=gmm)

#####
# Original fit.

# Get the optimal test set regressions.
reg_params_test = regs_test.get_optimal_regression_params()

# Get the test error for the original fit.
orig_test_error = mse_utils.get_lo_err_folded(
    opt_comb_params,
    keep_inds=full_lo_inds,
    mse_regs=regs_test,
    mse_reg_params=reg_params_test,
    gmm=gmm)

orig_pred = mse_utils.get_predictions(
    gmm, opt_comb_params['mix'], reg_params_test)

# Get the test error for the CV refit.
cv_error = mse_utils.get_lo_err_folded(
    comb_params_refit,
    keep_inds=full_lo_inds,
```

```

mse_regs=regs_test,
mse_reg_params=reg_params_test,
gmm=gmm)

cv_pred = mse_utils.get_predictions(
    gmm, comb_params_refit['mix'], reg_params_test)

# Get the test error for the IJ approximation.
ij_error = mse_utils.get_lo_err_folded(
    comb_params_lin,
    keep_inds=full_lo_inds,
    mse_regs=regs_test,
    mse_reg_params=reg_params_test,
    gmm=gmm)

ij_pred = mse_utils.get_predictions(
    gmm, comb_params_lin['mix'], reg_params_test)

```

Calculating prediction error.

1.0.3 Selected results.

We now make a cursory comparison of the results. For a more detailed analysis, including the results that went into the paper, see the notebook `examine_and_save_results`.

```

In [9]: cv_excess_error = cv_error - orig_test_error
        ij_excess_error = ij_error - orig_test_error

def GetColDf(col):
    return pd.DataFrame(
        {'cv_error': cv_error[:, col],
         'cv_excess': cv_excess_error[:, col],
         'ij_error': ij_error[:, col],
         'ij_excess': ij_excess_error[:, col],
         'col': col})

result = pd.concat([ GetColDf(col) for col in range(len(full_lo_inds)) ])


```

If we simply look at the point-by-point error, CV and IJ are highly correlated.

```

In [10]: sns.jointplot(x='cv_error', y='ij_error', data=result);


```



However, this is because the error in each point is dominated by the error at the original optimum. To meaningfully compare the IJ to CV, we should compare the difference between the IJ and CV error and the error at the original optimum. The distribution of these “difference-in-difference” errors is shown in the next plot.

Some clear outliers can be seen. However, note that, in this case, overplotting makes IJ look worse than it is – in the histograms you can see that most differences are very small.

```
In [11]: sns.jointplot(x='cv_excess', y='ij_excess', data=result);
```




As you might expect from a linear approximation, the IJ does the worst when the predicted change for CV is large.

```
In [12]: misfit = np.max(np.abs(cv_excess_error - ij_excess_error), axis=1)
         abs_cv_excess_error = np.max(np.abs(cv_excess_error), axis=1)

         sns.jointplot(abs_cv_excess_error, misfit)
```

```
Out[12]: <seaborn.axisgrid.JointGrid at 0x7f3f74fe1908>
```



Finally, we visualize some of the genes where IJ badly misestimates the CV error. Clearly, in these cases, re-fitting with the left-out points (shown with large dots) produced large changes that the IJ did not capture. In general, it appears that the IJ errs relative to CV by not moving far enough from the original optimum.

Despite the poor fit on these extreme genes, we stress that most genes exhibited small changes in both CV and IJ. For these genes, IJ performs well enough to capture salient aspects of the estimated out-of-sample error. For more detailed analysis of this point, see the notebook `examine_and_save_results`.

```
In [13]: timepoints = initial_metadata['timepoints']
         timepoints_stretch = np.sqrt(timepoints)

         def PlotGenePredictions(gene_ind):
             _, figs = plt.subplots(1, 3, figsize=(15,6))
```

```

for i in range(3):
    np.random.seed(42)
    plot_utils_lib.PlotRegressionLine(
        timepoints_stretch, regs_test, reg_params_test, gene_ind, this_plot=figs[i]
    )
    figs[i].plot(timepoints_stretch[full_lo_inds],
                 regs_test.y[gene_ind, full_lo_inds], 'o', markersize=10)

    plot_utils_lib.PlotPredictionLine(
        timepoints_stretch, regs_test, orig_pred, gene_ind, this_plot=figs[0])
    figs[0].set_title('Gene {} original fit'.format(gene_ind))


    plot_utils_lib.PlotPredictionLine(
        timepoints_stretch, regs_test, ij_pred, gene_ind, this_plot=figs[1])
    figs[1].set_title('Gene {} IJ fit'.format(gene_ind))

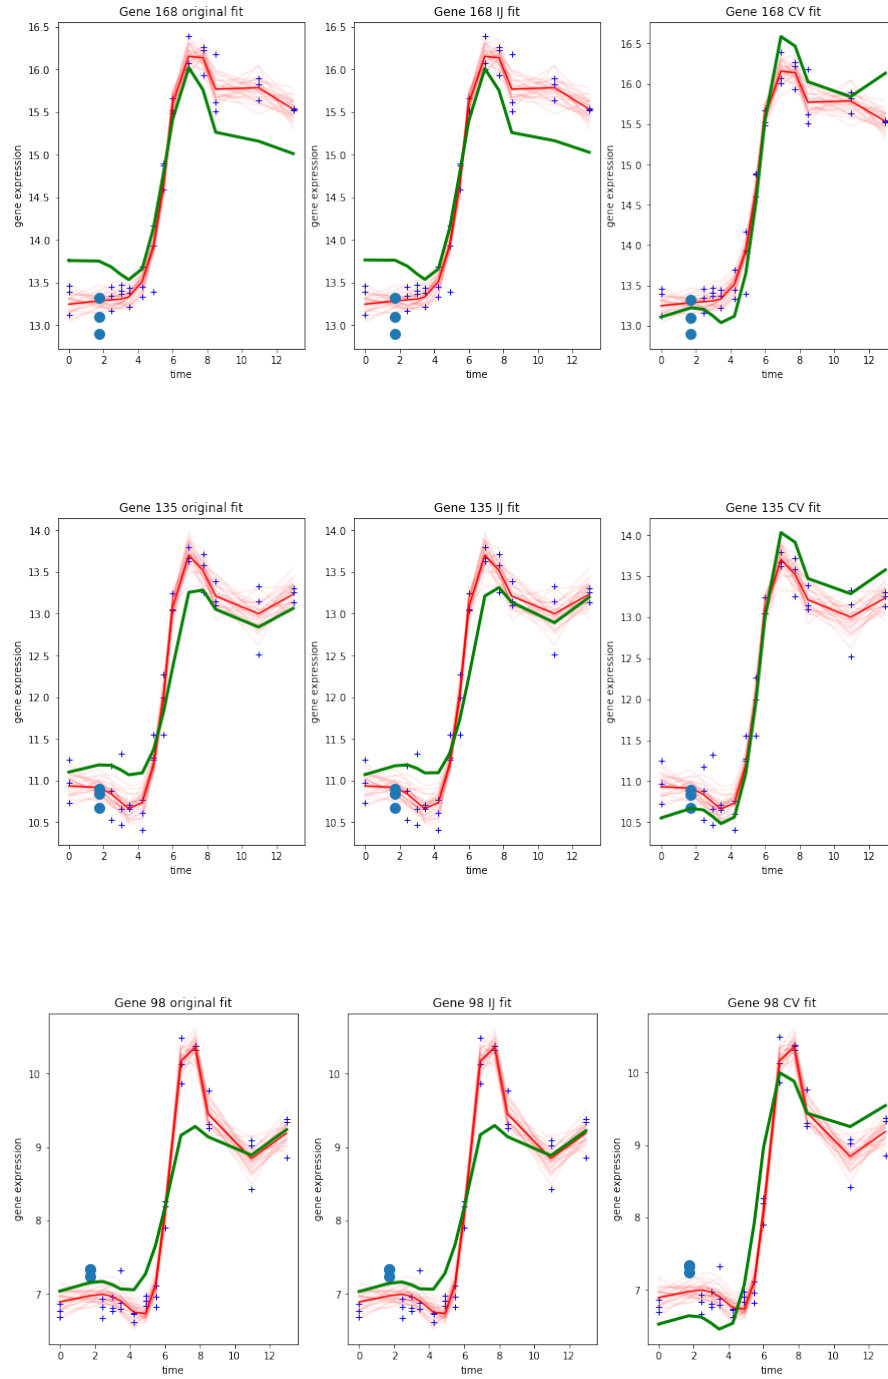
    plot_utils_lib.PlotPredictionLine(
        timepoints_stretch, regs_test, cv_pred, gene_ind, this_plot=figs[2])
    figs[2].set_title('Gene {} CV fit'.format(gene_ind))


In [14]: worst_fits = np.argsort(-1 * misfit)

for gene in worst_fits[0:5]:
    PlotGenePredictions(gene)

```







examine_and_save_results

February 21, 2019

1 Detailed analysis of results.

This notebook loads the output of the scripts in the directory `cluster_scripts` (particularly, the final script, `run_slurm_pred_error.py`). It produces the Rdata file that is used for the graphs in the paper as well as a number of supplemental analyses.

```
In [1]: library(tidyverse)
        library(gridExtra)
        library(repr) # For setting plot sizes
        source("load_python_data_lib.R")
        py_main <- InitializePython()
```

```
Attaching packages: tidyverse 1.2.1
  ggplot2 3.1.0      purrr  0.2.5
  tibble  1.4.2      dplyr  0.7.8
  tidyr   0.8.1      stringr 1.3.1
  readr   1.1.1      forcats 0.3.0
Conflicts: tidyverse_conflicts()
dplyr::filter() masks stats::filter()
dplyr::lag()    masks stats::lag()
```

Attaching package: gridExtra

The following object is masked from package:dplyr:

combine

Attaching package: reshape2

The following object is masked from package:tidyr:

smiths

```
In [2]: # Choose the initialization method.
        init_method <- "kmeans" # This is the choice for the paper.
```

```

#init_method <- "warm"

# Choose whether or not to re-run the regressions before calculating test error.
use_rereg <- FALSE # This is the choice for the paper.
#use_rereg <- TRUE

# This is the file that is used in the paper's knitr.
save_dir <- "../..fits"
save_filename <- sprintf("paper_results_init_%s_rereg_%s.Rdata", init_method, use_rereg)

```

1.0.1 Load the saved data for all dfs and k

```

In [3]: dfs <- list()
        metadata_dfs <- list()

        for (lo_num_times in 1:3) {
          cat("lo_num_times ", lo_num_times)
          for (df in 4:8) {
            cat(".")
            load_res <- LoadPredictionError(df, lo_num_times, init_method)
            this_refit_err_df <- load_res$refit_err_df
            this_metadata_df <- load_res$metadata_df
            this_refit_err_melt <- MeltErrorColumns(this_refit_err_df)
            dfs[[length(dfs) + 1]] <- this_refit_err_melt
            metadata_dfs[[length(metadata_dfs) + 1]] <- this_metadata_df
          }
          cat("\n")
        }
        cat("Done.\n")
        refit_err_melt <- do.call(bind_rows, dfs)
        metadata_df <- do.call(bind_rows, metadata_dfs)

lo_num_times 1...
lo_num_times 2...
lo_num_times 3...
Done.

```

1.0.2 Metadata (timing, parameter dimensions)

Make a tidy dataframe with the metadata. The parameter length, Hessian time, and initial optimization time are all reported in the text of the paper. Their values will be derived from this dataframe in knitr.

```

In [4]: metadata_df <-
        metadata_df %>%
        mutate(lr_hess_time=total_lr_time + initial_hess_time,
               avg_lr_time=total_lr_time / num_comb,
               avg_refit_time=total_refit_time / num_comb,

```

```

        param_length=gmm_param_length + reg_param_length)
print(names(metadata_df))

select(metadata_df, df, param_length) %>%
  group_by(df) %>%
  summarize(param_length=unique(param_length))

select(metadata_df, df, initial_hess_time, initial_opt_time) %>%
  group_by(df) %>%
  summarize(initial_hess_time=median(initial_hess_time),
            initial_opt_time=median(initial_opt_time))

round(median(metadata_df$initial_opt_time), digits=-1)

[1] "num_comb"          "total_lr_time"      "total_refit_time"
[4] "initial_opt_time"  "initial_reg_time"   "initial_hess_time"
[7] "gmm_param_length"  "reg_param_length"   "df"
[10] "lo_num_times"      "init_method"        "lr_hess_time"
[13] "avg_lr_time"       "avg_refit_time"     "param_length"

```

df	param_length
----	--------------

4	25325
5	31643
6	38661
7	46379
8	54797

df	initial_hess_time	initial_opt_time
----	-------------------	------------------

4	275.7295	31.44656
5	295.0325	41.84182
6	359.6855	35.11145
7	478.7345	50.88843
8	584.4987	77.02919

40

Make a dataframe for the timing plot from the metadata.


```

In [5]: metadata_graph_df <-
        metadata_df %>%
        select(df, lo_num_times, total_refit_time, lr_hess_time, initial_opt_time) %>%
        melt(id.vars=c("lo_num_times", "df"))
        head(metadata_graph_df)

```

lo_num_times	df	variable	value
1	4	total_refit_time	338.1638
1	5	total_refit_time	391.6006
1	6	total_refit_time	423.8322
1	7	total_refit_time	632.2635
1	8	total_refit_time	599.0894
2	4	total_refit_time	1123.7316

```
In [6]: ggplot(metadata_graph_df) +
  geom_bar(aes(x=lo_num_times, y=value, fill=variable, group=variable),
    stat="identity", position=position_dodge()) +
  facet_grid( ~ df)
```



1.0.3 Calculate prediction errors

Make summaries of prediction error for various methods and datasets.

```
In [7]: # In-sample IJ error.
lr_df <-
  refit_err_melt %>%
  filter(rereg==use_rereg, method=="lin", test==FALSE, measure=="err") %>%
```

```

  rename(error=value) %>%
  mutate(output="lin_in_sample")

# In-sample CV error.
cv_df <-
  refit_err_melt %>%
  filter(rereg==use_rereg, method=="ref", test==FALSE, measure=="err") %>%
  rename(error=value) %>%
  mutate(output="cv_in_sample")

# In-sample training error (no points left out).
train_df <-
  refit_err_melt %>%
  filter(rereg==use_rereg, method=="ref", test==FALSE, measure=="train_err") %>%
  rename(error=value) %>%
  mutate(output="train_error")

# Out-of-sample test error.
test_df <-
  refit_err_melt %>%
  filter(rereg==use_rereg, method=="ref", test==TRUE, measure=="train_err") %>%
  rename(error=value) %>%
  mutate(output="test_error")

refit_for_df_choice <- bind_rows(
  lr_df, cv_df, test_df, train_df)

```

In [8]: `head(refit_for_df_choice)`

test	method	comb	rereg	gene	df	lo_num_times	time	measure	error	output
FALSE	lin	0	FALSE	0	4	1	0	err	1.0088933	lin_in_sam
FALSE	lin	0	FALSE	1	4	1	0	err	0.1243607	lin_in_sam
FALSE	lin	0	FALSE	2	4	1	0	err	-0.4340983	lin_in_sam
FALSE	lin	0	FALSE	3	4	1	0	err	-0.2203431	lin_in_sam
FALSE	lin	0	FALSE	4	4	1	0	err	1.9032786	lin_in_sam
FALSE	lin	0	FALSE	5	4	1	0	err	-0.2876837	lin_in_sam

Make a tidy dataframe for choosing df. The graph in the paper will be based on this dataframe.

Note that most of the signal for choosing df is already in the training data error. However, there is an uptick in error in both CV and IJ for df=8 which is not captured by the training data error.

```

In [9]: refit_err_summary <-
  refit_for_df_choice %>%
  group_by(output, df, lo_num_times) %>%
  mutate(esize=abs(error)) %>%
  summarize(med=median(esize),
            mean=mean(esize),
            n_obs=n(),
            se=sd(esize) / sqrt(n_obs),


```

```

qlow=quantile(esize, 0.25),
qhigh=quantile(esize, 0.75))

ggplot(refit_err_summary) +
  geom_line(aes(x=df, y=mean, group=output, color=output)) +
  geom_errorbar(aes(x=df, ymin=mean - 2 * se, ymax=mean + 2 * se,
                    group=output, color=output)) +
  facet_grid(~ lo_num_times) +
  ggtitle(sprintf("%d times left out", lo_num_times))

```



1.0.4 Gene-by-gene accuracy measures.

```

In [10]: refit_err_plot <-
         refit_err_melt %>%

```

```

filter(rereg==use_rereg) %>%
dcast(df + lo_num_times + test + comb + rereg + gene + time ~ method + measure,
      value.var=c("value"))

```

We now look at the correlation between the CV and IJ prediction errors across genes. For each df and k , there are a number of different combinations of left-out points. We report the median, min, and max correlation coefficients across these combinations of left-out points.

First, we show the correlation between the raw prediction errors. Although the correlation is quite high, this is because the training error at the original optimum is the principle source of variation in the errors across genes, and this quantity is common to both CV and IJ.

```

In [11]: err_corr <- refit_err_plot %>%
        filter(test==FALSE, rereg==use_rereg) %>%
        group_by(df, lo_num_times, comb) %>%
        summarize(r=cor(lin_err, ref_err)) %>%
        group_by(df, lo_num_times) %>%
        summarize(med_r=median(r), min_r=min(r), max_r=max(r))

        print("Correlation between error: ")
        print(err_corr)

```

```

[1] "Correlation between error: "
# A tibble: 15 x 5
# Groups:   df [?]
   df lo_num_times med_r min_r max_r
<int> <int> <dbl> <dbl> <dbl>
1     4           1 0.974 0.949 0.984
2     4           2 0.975 0.902 0.992
3     4           3 0.967 0.871 0.991
4     5           1 0.963 0.856 0.983
5     5           2 0.966 0.860 0.984
6     5           3 0.947 0.759 0.981
7     6           1 0.980 0.807 0.985
8     6           2 0.968 0.835 0.986
9     6           3 0.929 0.759 0.983
10    7           1 0.962 0.794 0.974
11    7           2 0.952 0.737 0.976
12    7           3 0.914 0.599 0.974
13    8           1 0.962 0.703 0.971
14    8           2 0.941 0.663 0.974
15    8           3 0.829 0.251 0.958

```

A more meaningful measure is the correlation in the excess error for IJ and CV over the error at the original fit.

```

In [12]: diff_corr <- refit_err_plot %>%
        filter(test==FALSE, rereg==use_rereg) %>%
        group_by(df, lo_num_times, comb) %>%

```

```

summarize(r=cor(lin_e_diff, ref_e_diff)) %>%
group_by(df, lo_num_times) %>%
summarize(med_r=median(r), min_r=min(r), max_r=max(r))

print("Correlation between difference from train error: ")
print(diff_corr)

[1] "Correlation between difference from train error: "
# A tibble: 15 x 5
# Groups:   df [?]
      df lo_num_times med_r   min_r max_r
  <int>   <int> <dbl>   <dbl> <dbl>
1     4         1  0.483  0.0956 0.844
2     4         2  0.577  0.277  0.828
3     4         3  0.605  0.303  0.833
4     5         1  0.464  0.143  0.728
5     5         2  0.510  0.330  0.709
6     5         3  0.510  0.312  0.671
7     6         1  0.655  0.368  0.783
8     6         2  0.588  0.218  0.845
9     6         3  0.499  0.0701 0.737
10    7         1  0.660  0.512  0.760
11    7         2  0.564  0.224  0.863
12    7         3  0.491  0.0344 0.801
13    8         1  0.744  0.380  0.900
14    8         2  0.646  0.166  0.862
15    8         3  0.214 -0.226  0.767


```

For higher degrees of freedom, increasing the number of left-out points seems to decrease the IJ's accuracy, as you might expect.

```

In [13]: ggplot(diff_corr) +
  geom_bar(aes(x=paste(df, lo_num_times, sep=","),
    y=med_r, fill=as.character(df)), stat="identity")

```

Plot the densities of the IJ and CV with points to show outliers. This is a graphical version of the results summarized by the correlation tables above.

```
In [14]: # There are a few outliers, so limit the extent of the plot so that
# the bulk of the distribution is visible.
qlim <- quantile(refit_err_plot$ref_e_diff, c(0.1, 0.9))

options(repr.plot.width=4, repr.plot.height=20)


# This plot, or ones like it, is probably the best measure of
# the accuracy of the IJ.
ggplot(filter(refit_err_plot, test == FALSE, lo_num_times==1)) +
  geom_point(aes(x=ref_e_diff, y=lin_e_diff), alpha=0.01) +
  geom_density2d(aes(x=ref_e_diff, y=lin_e_diff)) +
```

```
geom_abline(aes(slope=1, intercept=0)) +  
facet_grid(df ~ rereg) +  
xlim(qlim[1], qlim[2]) + ylim(qlim[1], qlim[2])
```

Warning message:

Removed 10770 rows containing non-finite values (stat_density2d).Warning message:

Removed 10770 rows containing missing values (geom_point).



1.0.5 Save results for plotting in the paper.

```
In [15]: print(sprintf("Saving to %s", file.path(save_dir, save_filename)))
        save(refit_err_summary,
              metadata_df,
              diff_corr,
              err_corr,
              file=file.path(save_dir, save_filename))

[1] "Saving to ../../fits/paper_results_init_kmeans_rereg_FALSE.Rdata"
```